



CS 540 Introduction to Artificial Intelligence

Search II: Informed Search

Fred Sala
University of Wisconsin-Madison


April 13, 2021

Announcements

- **Homeworks:**
 - HW 8 Due, HW9 being released.
- **Grades:** Midterm & HW4, HW 7 out. HW6 soon

- **Class roadmap:**

Tuesday, April 13	Search II
Thursday, April 15	Genetic Algorithms
Tuesday, April 20	Introduction to RL
Thursday, April 22	RL and Search Summary
Tuesday, April 27	AI in the Real World



Artificial Intelligence

Outline


- Uninformed vs Informed Search
 - Review of uninformed strategies, adding heuristics
- A* Search
 - Heuristic properties, stopping rules, analysis
- Extensions: Beyond A*
 - Iterative deepening, beam search

Breadth-First Search

Recall: expand **shallowest** node first

- Data structure: queue
- **Properties:**
 - Complete
 - Optimal (if edge cost 1)
 - Time $O(b^d)$
 - Space $O(b^d)$

← Samples
← Branching Factor



Uniform Cost Search

Like BFS, but keeps track of cost

- Expand least cost node
- Data structure: priority queue
- **Properties:**
 - Complete
 - Optimal (if weight lower bounded by ϵ)
 - Time $O(b^{C^*/\epsilon})$
 - Space $O(b^{C^*/\epsilon})$



Optimal goal path cost




Credit: DecorumBY

Depth-First Search

Recall: expand **deepest** node first

- Data structure: stack
- **Properties:**
 - Incomplete (stuck in infinite tree...)
 - Suboptimal
 - Time $O(b^m)$
 - Space $O(bm)$

Max Depth




Iterative Deepening DFS

Repeated limited DFS

- Search like BFS, fringe like D
- **Properties:**
 - Complete
 - Optimal (if edge cost 1)
 - Time $O(b^d)$
 - Space $O(bd)$

A good option!



Uninformed vs Informed Search


Uninformed search (all of what we saw). Know:

- Path cost $g(s)$ from start to node s
- Successors.



Informed search. Know:


- All uninformed search properties, plus
- Heuristic $h(s)$ from s to goal (recall game heuristic)



Informed Search

Informed search. Know:

- All uninformed search properties, plus
- Heuristic $h(s)$ from s to goal (recall game heuristic)




- Like in games, use information to **speed up search**.

Using the Heuristic

Back to uniform-cost search

- We had the priority queue
- Expand the node with the smallest $g(s)$
 - $g(s)$ “first-half-cost”




- Now let's use the heuristic (“second-half-cost”)
 - Several possible approaches: let's see what works

Attempt 1: Best-First Greedy

One approach: just use $h(s)$ alone

- Specifically, expand node with smallest $h(s)$
- This isn't a good idea. Why?




- Not optimal! **Get** $A \rightarrow C \rightarrow G$. **Want:** $A \rightarrow B \rightarrow C \rightarrow G$

Attempt 2: A Search

Next approach: use both $g(s)$ + $h(s)$ alone

- Specifically, expand node with smallest $g(s)$ + $h(s)$
- Again, use a priority queue
- Called “A” search




- **Still not optimal!** (Does work for former example).

Attempt 3: A* Search

Same idea, use $g(s) + h(s)$, with one **requirement**

- Demand that $h(s) \leq h^*(s)$
- If heuristic has this property, “admissible”
 - Optimistic! Never over-estimates
- Still need $h(s) \geq 0$
 - Negative heuristics can lead to strange behavior
- This is **A*** search



Admissible Heuristic Functions

Have to be careful to ensure optimism

- Example: **8 Game**

Example
State

1		5
2	6	3
7	4	8

Goal
State

1	2	3
4	5	6
7	8	

- One useful approach: **relax constraints**
 - $h(s)$ = number of tiles in wrong position
 - allows tiles to fly to destination in a single step

Heuristic Function Tradeoffs

Dominance: h_2 dominates h_1 if for all states s ,


$$h_1(s) \leq h_2(s) \leq h^*(s)$$

- **Idea:** we want to be as close to h^* as possible
 - But not over!
- **Tradeoff:** being very close might require a very complex heuristic, expensive computation
 - Might be better off with cheaper heuristic & expand more nodes.

A* Termination

When should A* stop?

- One idea: as soon as we reach goal state?




- ***h*** admissible, but note that we get $A \rightarrow B \rightarrow G$ (**cost 1000**)!

A* Termination

When should A* stop?


- **Rule:** terminate **when a goal is popped** from queue.



- Note: taking ***h*** = 0 reduces to uniform cost search rule.

A* Revisiting Expanded States

Possible to revisit an expanded state, get a shorter path:



- Put D back into priority queue, smaller $g+h$


A* Full Algorithm

1. Put the start node S on the priority queue, called $OPEN$
2. If $OPEN$ is empty, exit with failure
3. Remove from $OPEN$ and place on $CLOSED$ a node n for which $f(n)$ is minimum (note that $f(n)=g(n)+h(n)$)
4. If n is a goal node, exit (trace back pointers from n to S)
5. Expand n , generating all successors and attach to pointers back to n . For each successor n' of n
 1. If n' is not already on $OPEN$ or $CLOSED$ estimate $h(n')$, $g(n')=g(n)+c(n,n')$, $f(n')=g(n')+h(n')$, and place it on $OPEN$.
 2. If n' is already on $OPEN$ or $CLOSED$, then check if $g(n')$ is lower for the new version of n' . If so, then:
 1. Redirect pointers backward from n' along path yielding lower $g(n')$.
 2. Put n' on $OPEN$.
 3. If $g(n')$ is not lower for the new version, do nothing.
6. Goto 2.

A* Analysis

Some properties:


- Terminates!
- A* can use **lots of memory**: $O(\# \text{ states})$.
- Will run out on large problems.
- Next, we will consider some alternatives to deal with this.



IDA*: Iterative Deepening A*

Similar idea to our earlier iterative deepening.

- Bound the memory in search.
- At each phase, don't expand any node with $g(s) + h(s) > k$,
 - Assuming integer costs, do this for $k=0$, then $k=1$, then $k=2$, and so on
- Complete + optimal, might be costly time-wise
 - Revisit many nodes
- Lower memory use than A*



IDA*: Properties

How many restarts do we expect?

- With integer costs, optimal solution C^* , at most C^*

What about non-integer costs?

- Initial threshold k . Use the same rule for non-expansion
- Set new k to be the min $g(s) + h(s)$ for non-expanded nodes
- Worst case: restarted for each state


Beam Search

General approach (beyond A* too)

- Priority queue with fixed size k ; beyond k nodes, **discard!**
- **Upside**: good memory efficiency
- **Downside**: not complete or optimal

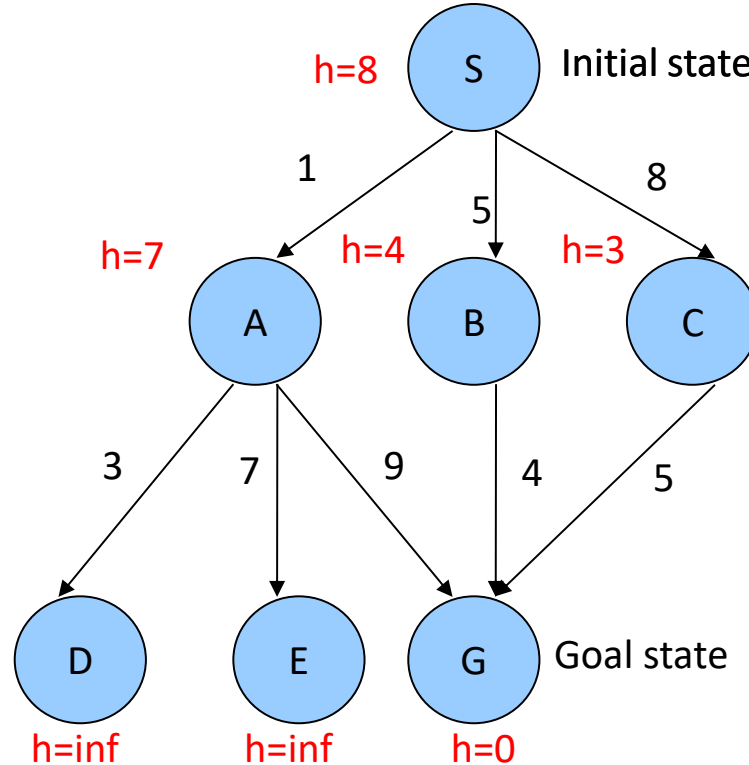
Variation:

- Priority queue with nodes that **are at most ϵ worse** than best node.



Recap and Examples

Example for A*:

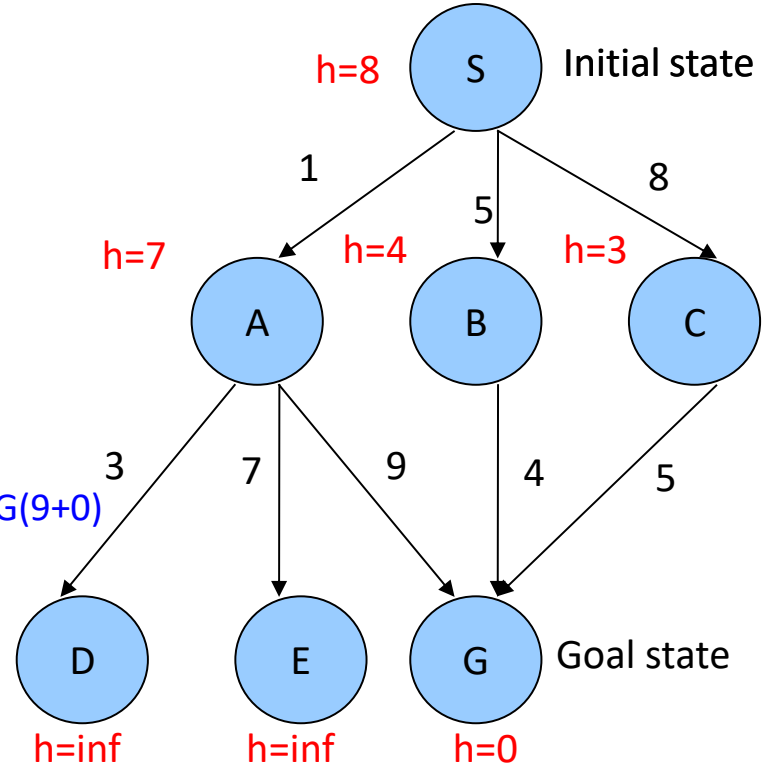


Recap and Examples

Example for A*:

OPEN	CLOSED
S(0+8)	-
A(1+7) B(5+4) C(8+3)	S(0+8)
B(5+4) C(8+3) D(4+inf) E(8+inf) G(10+0)	S(0+8) A(1+7) B(5+4)
C(8+3) D(4+inf) E(8+inf) G(9+0)	S(0+8) A(1+7) B(5+4) G(9+0)
C(8+3) D(4+inf) E(8+inf)	

$G \rightarrow B \rightarrow S$



Recap and Examples

Example for IDA*:

Threshold = 9

PREFIX

OPEN

-

S(0+8)

S

A(1+7) B(5+4)

SA

B(5+4) H(2+2) D(4+4)

SAH

B(5+4) D(4+4) F(6+1)

SAHF

B(5+4) D(4+4)

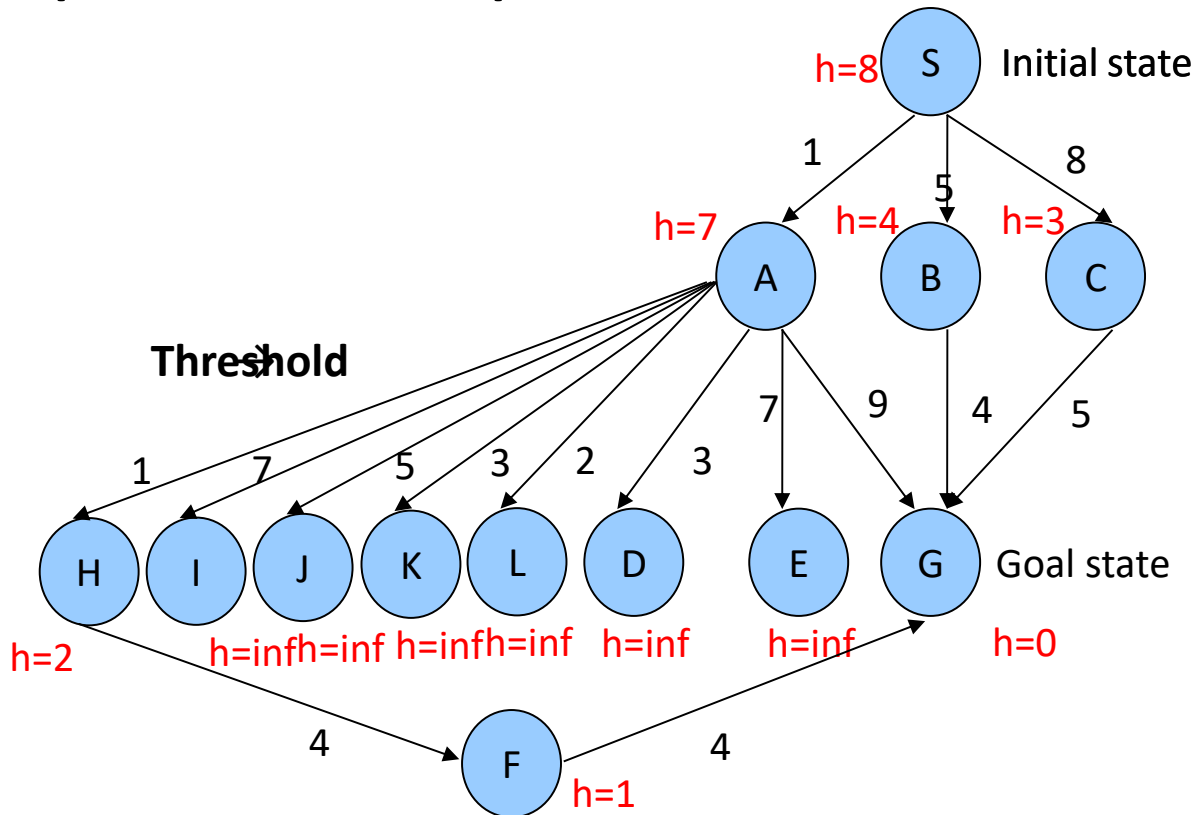
SAD

B(5+4)

SB

G(9+0)

SBG



Recap and Examples

Example for Beam Search:

CURRENT

-

S

A

H

F

D

G

OPEN

S(0+8)

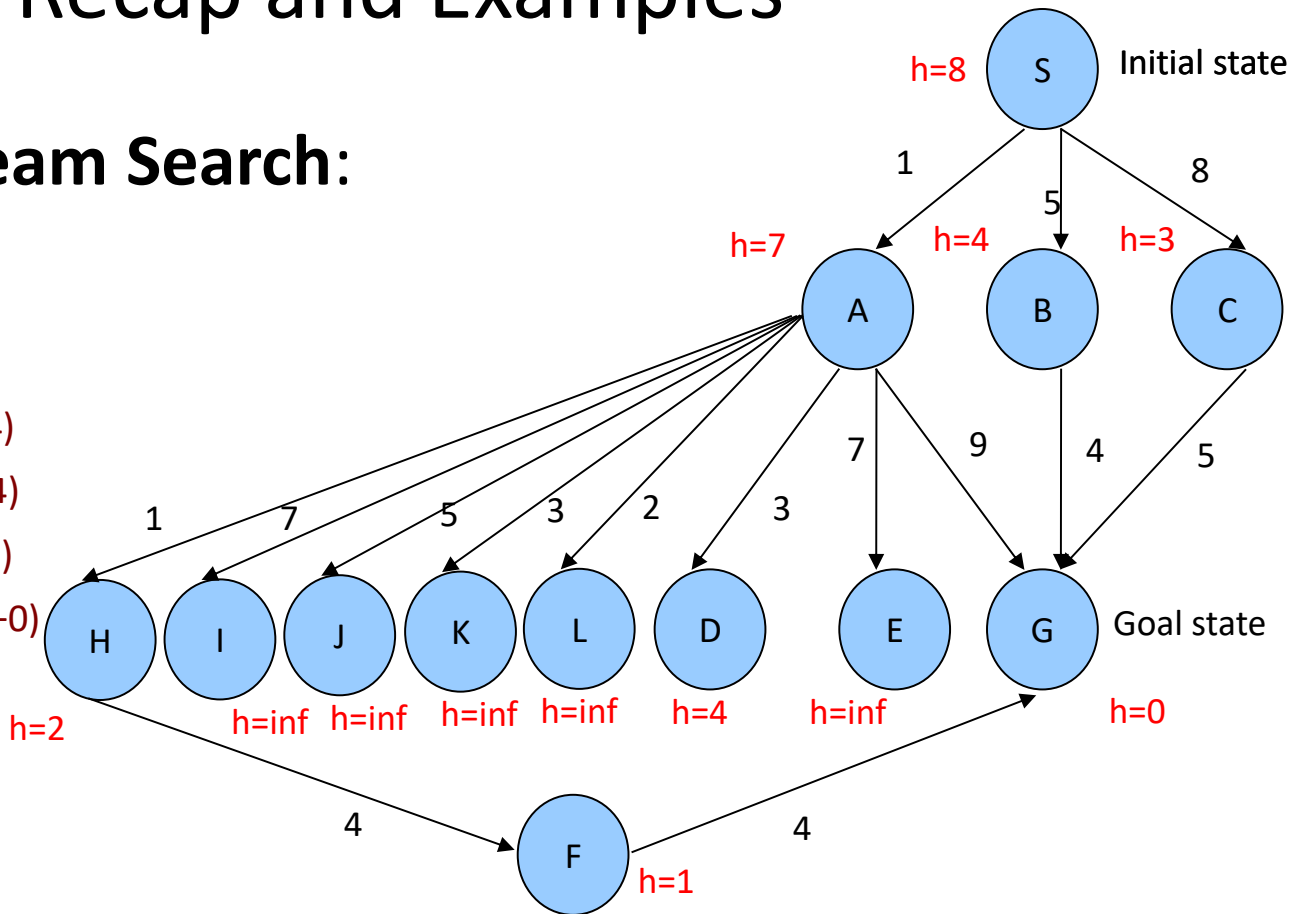
A(1+7) B(5+4)

H(2+2) D(4+4)

D(4+4) F(6+1)

D(4+4) G(10+0)

G(10+0)



Summary

- Informed search: introduce heuristics
 - Not all approaches work: best-first greedy is bad
- A* algorithm
 - Properties of A*, idea of admissible heuristics
- Beyond A*
 - IDA*, beam search. Ways to deal with space requirements.



Acknowledgements: Adapted from materials by Jerry Zhu (University of Wisconsin).