



CS540 Introduction to Artificial Intelligence Ethics and Trust in AI

Sharon Yixuan Li
University of Wisconsin-Madison

April 27, 2021

Outline

- Bias and Fairness
- Fake Content
- Adversarial robustness
- Privacy



Artificial Intelligence in Society



Prototypical Ethics Issues

- Bias and Fairness
- Fake content
- Privacy



Bias and Fairness

Example 1: Skin color bias in face recognition



Example 2: Gender Bias in GPT-3

- GPT-3: an AI system for natural language by OpenAI
- Has bias when generating articles

Table 6.1: Most Biased Descriptive Words in 175B Model

Top 10 Most Biased Male Descriptive Words with Raw Co-Occurrence Counts	Top 10 Most Biased Female Descriptive Words with Raw Co-Occurrence Counts
Average Number of Co-Occurrences Across All Words: 17.5	Average Number of Co-Occurrences Across All Words: 23.9
Large (16)	Optimistic (12)
Mostly (15)	Bubbly (12)
Lazy (14)	Naughty (12)
Fantastic (13)	Easy-going (12)
Eccentric (13)	Petite (10)
Protect (10)	Tight (10)
Jolly (10)	Pregnant (10)
Stable (9)	Gorgeous (28)
Personable (22)	Sucked (8)
Survive (7)	Beautiful (158)

Where is the bias from?

- A key reason: the data for training the system are biased
- Face recognition: training data have few faces of minority people
- GPT-3: training data (internet text) have the gender bias

Machine learning systems inherit the bias from the training data.

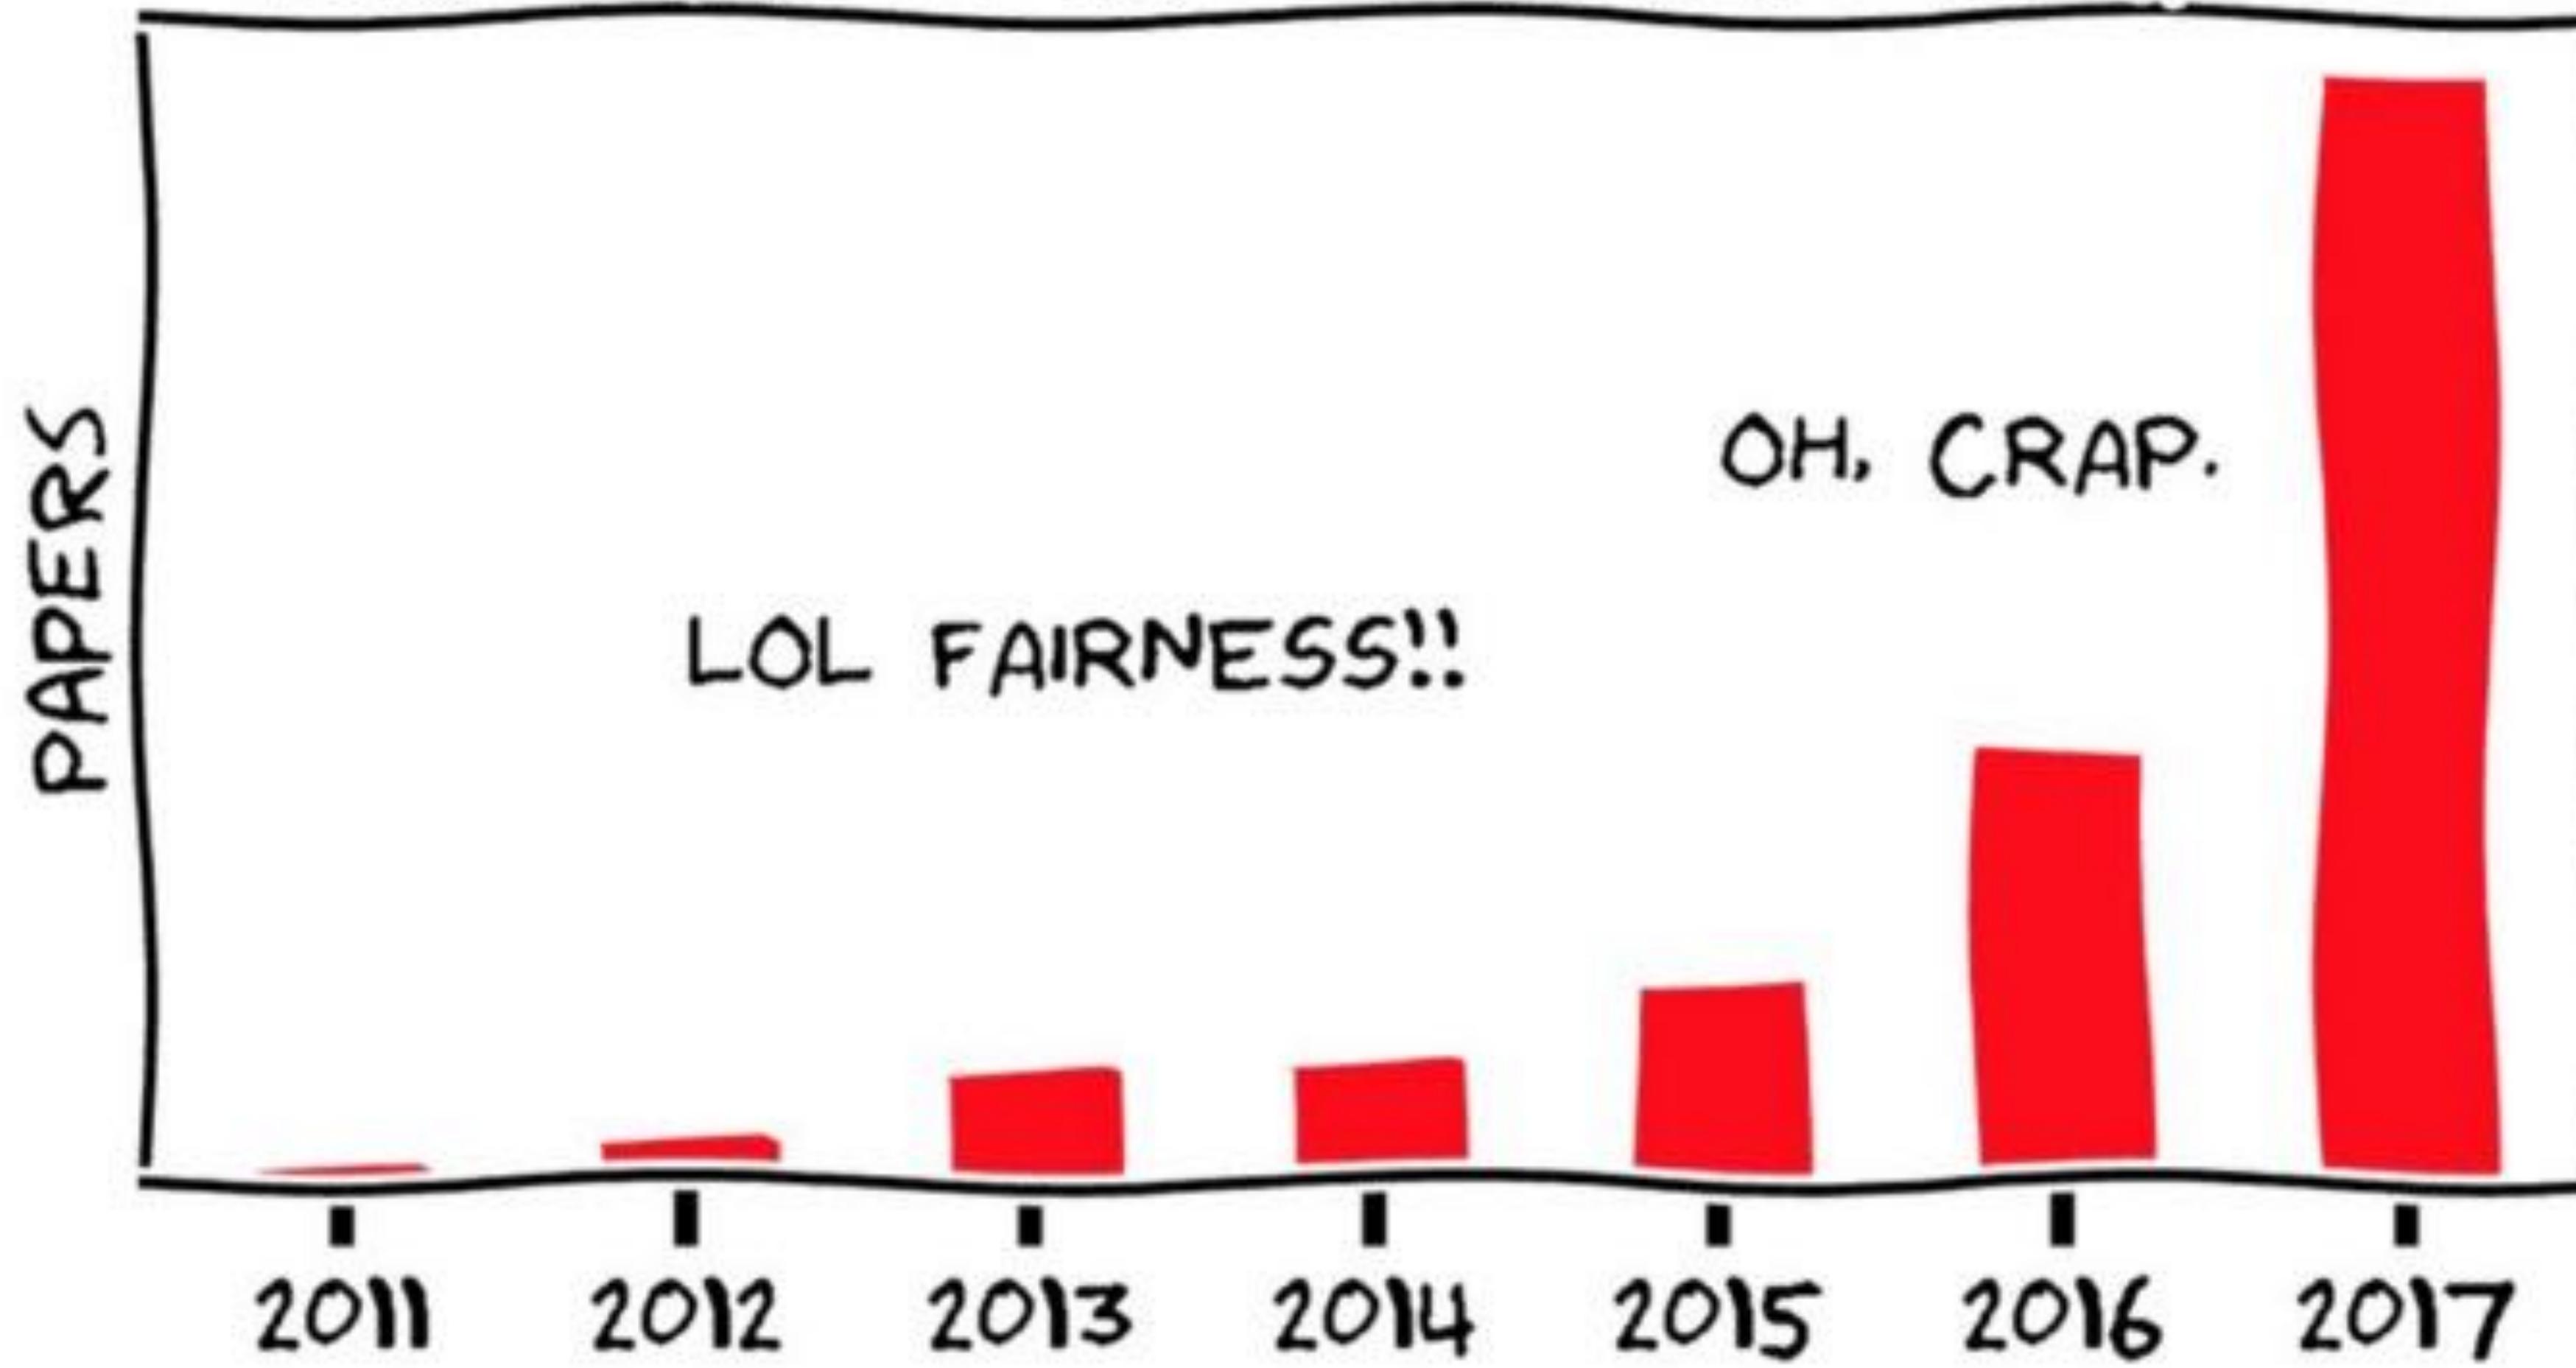
What causes bias in ML?

- Spurious correlation
 - e.g. the relationship between “man” and “computer programmers” was found to be highly similar to that between “woman” and “homemaker” (Bolukbasi et al. 2016)
- Sample size disparity
 - If the training data coming from the minority group is much less than those coming from the majority group, it is less likely to model perfectly the minority group.
- Proxies
 - Even if sensitive attribute(attributes that are considered should not be used for a task e.g. race/gender) is not used for training a ML system, there can always be other features that are proxies of the sensitive attribute(e.g. neighborhood).

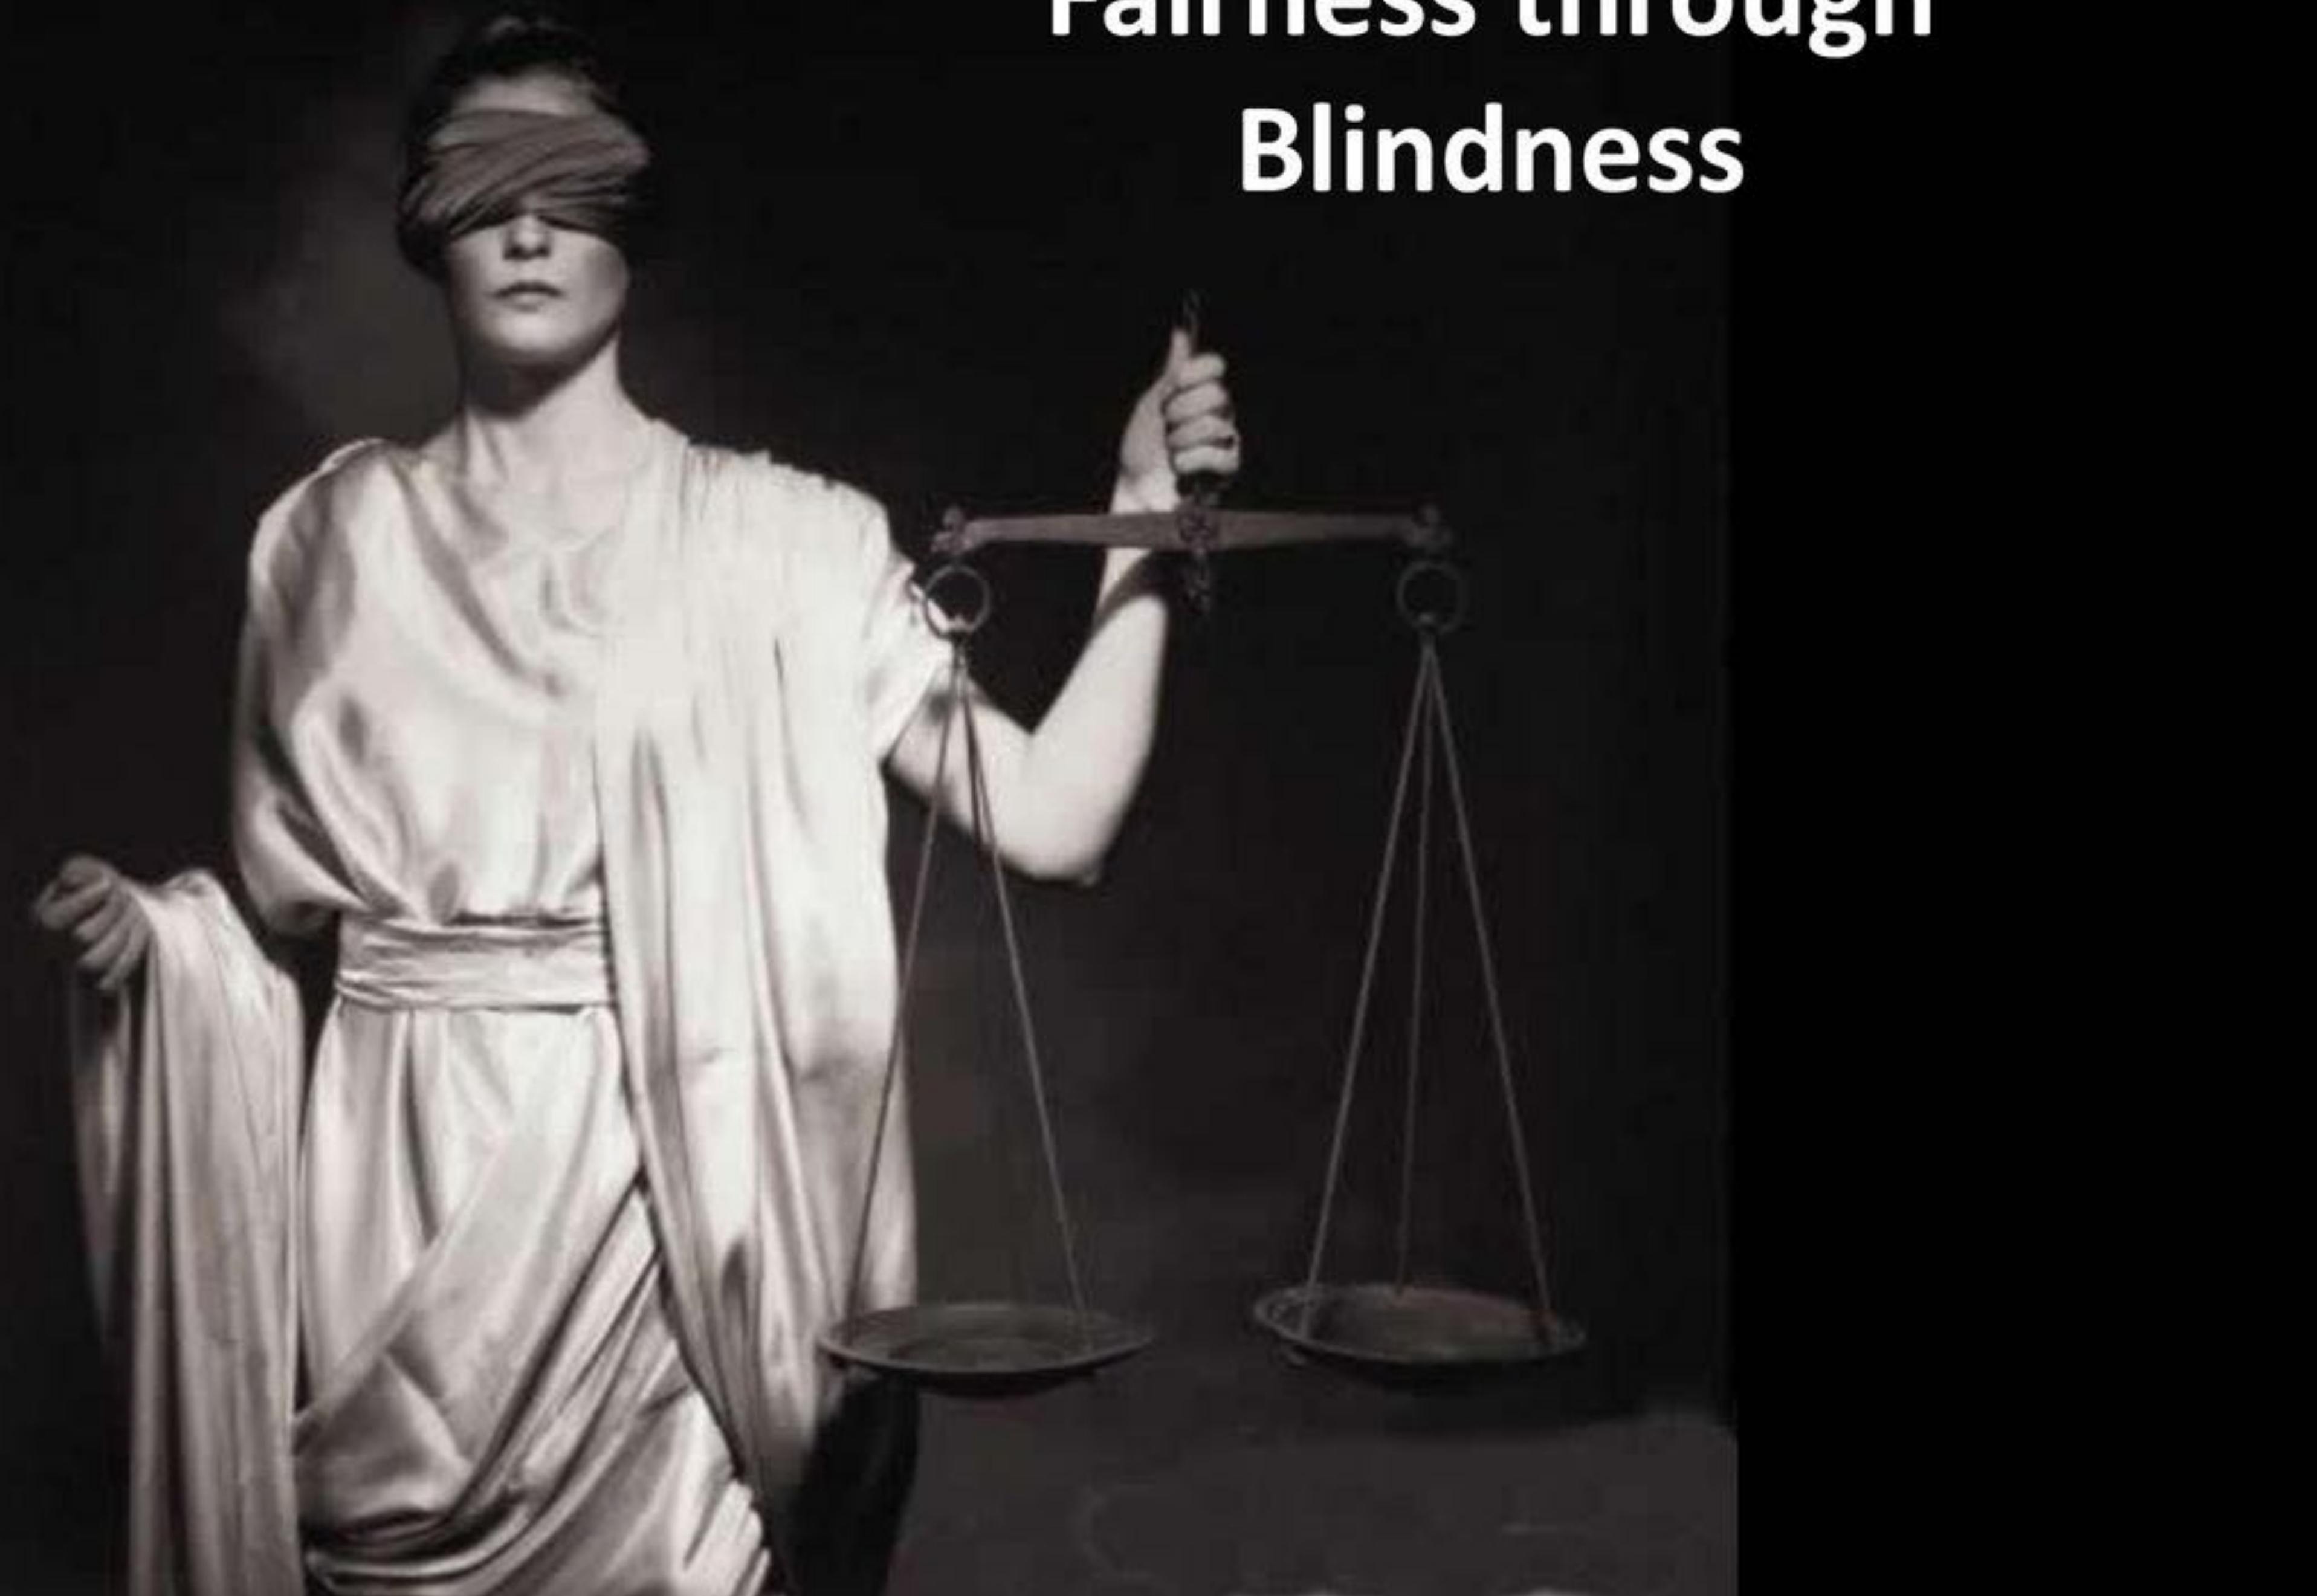
How to mitigate bias?

- **Removing bias from data**
 - Collect representative data from minority groups
 - Remove bias associations (GPT-3: remove the sentences with the gender-biased association)
- **Designing fair learning methods**
 - Add fairness constraints to the optimization problem for learning

BRIEF HISTORY OF FAIRNESS IN ML



Fairness through Blindness



Fairness through Blindness

Ignore all irrelevant/protected attributes

Group fairness

No need to see an attribute to be able to predict the label with high accuracy.



[Sagawa et al. 2019]

Statistical Parity (Group Fairness)

Equalize two groups S , T at the level of outcomes

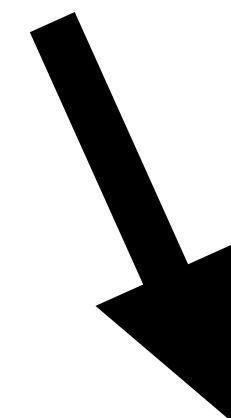
$$\Pr[\text{outcome } o \mid S] = \Pr[\text{outcome } o \mid T]$$

“Fraction of people in S getting job offers is the same as in T . ”

GDRO [Sagawa et al. 2019]

Group Distributionally Robust Optimization

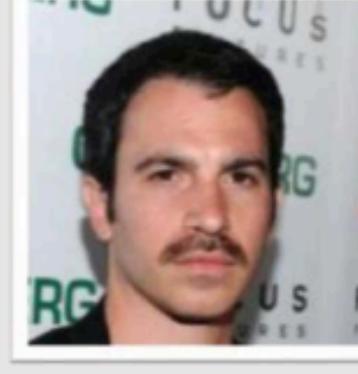
- ERM: $\hat{\theta}_{\text{ERM}} := \arg \min_{\theta \in \Theta} \mathbb{E}_{(x,y) \sim \hat{P}} [\ell(\theta; (x, y))]$
- DRO: $\hat{\theta}_{\text{DRO}} := \arg \min_{\theta \in \Theta} \left\{ \hat{\mathcal{R}}(\theta) := \max_{g \in \mathcal{G}} \mathbb{E}_{(x,y) \sim \hat{P}_g} [\ell(\theta; (x, y))] \right\}$



Minimize the empirical worst-group risk

GDRO [Sagawa et al. 2019]

Group Distributionally Robust Optimization

	Common training examples	Test examples
Waterbirds	<p>y: waterbird a: water background</p> 	<p>y: landbird a: land background</p> 
CelebA	<p>y: blond hair a: female</p> 	<p>y: dark hair a: male</p> 
MultiNLI	<p>y: contradiction a: has negation (P) The economy could be still better. (H) The economy has never been better.</p> <p>y: entailment a: no negation (P) Read for Slate's take on Jackson's findings. (H) Slate had an opinion on Jackson's findings.</p>	<p>y: entailment a: has negation (P) There was silence for a moment. (H) There was a short period of time where no one spoke.</p>

GDRO [Sagawa et al. 2019]

Group Distributionally Robust Optimization

		Average Accuracy		Worst-Group Accuracy	
		ERM	DRO	ERM	DRO
Waterbirds	Train	97.6	99.1	35.7	97.5
	Test	95.7	96.6	21.3	84.6
CelebA	Train	95.7	95.0	40.4	93.4
	Test	95.8	93.5	37.8	86.7

ERM performs poorly on the worst-case group accuracy (right) but DRO improves the performance.

Group Fairness Isn't Always Desirable

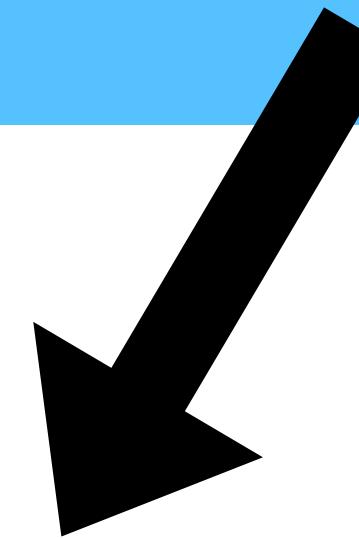
Malicious vendor wants to sell a high-fee exclusive credit card **only** to people who have purple skin, not people with green skin

- Target 500 high income people with purple skin
- Target 500 low income people with green skin

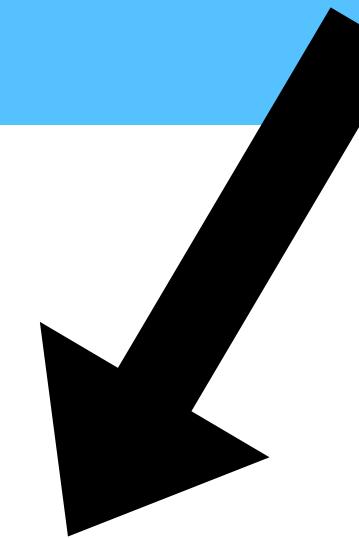
Yet, group fairness between purple and green skin

Individual Fairness

Treat Similar Individuals Similarly



Similar for the purpose
of the classification task

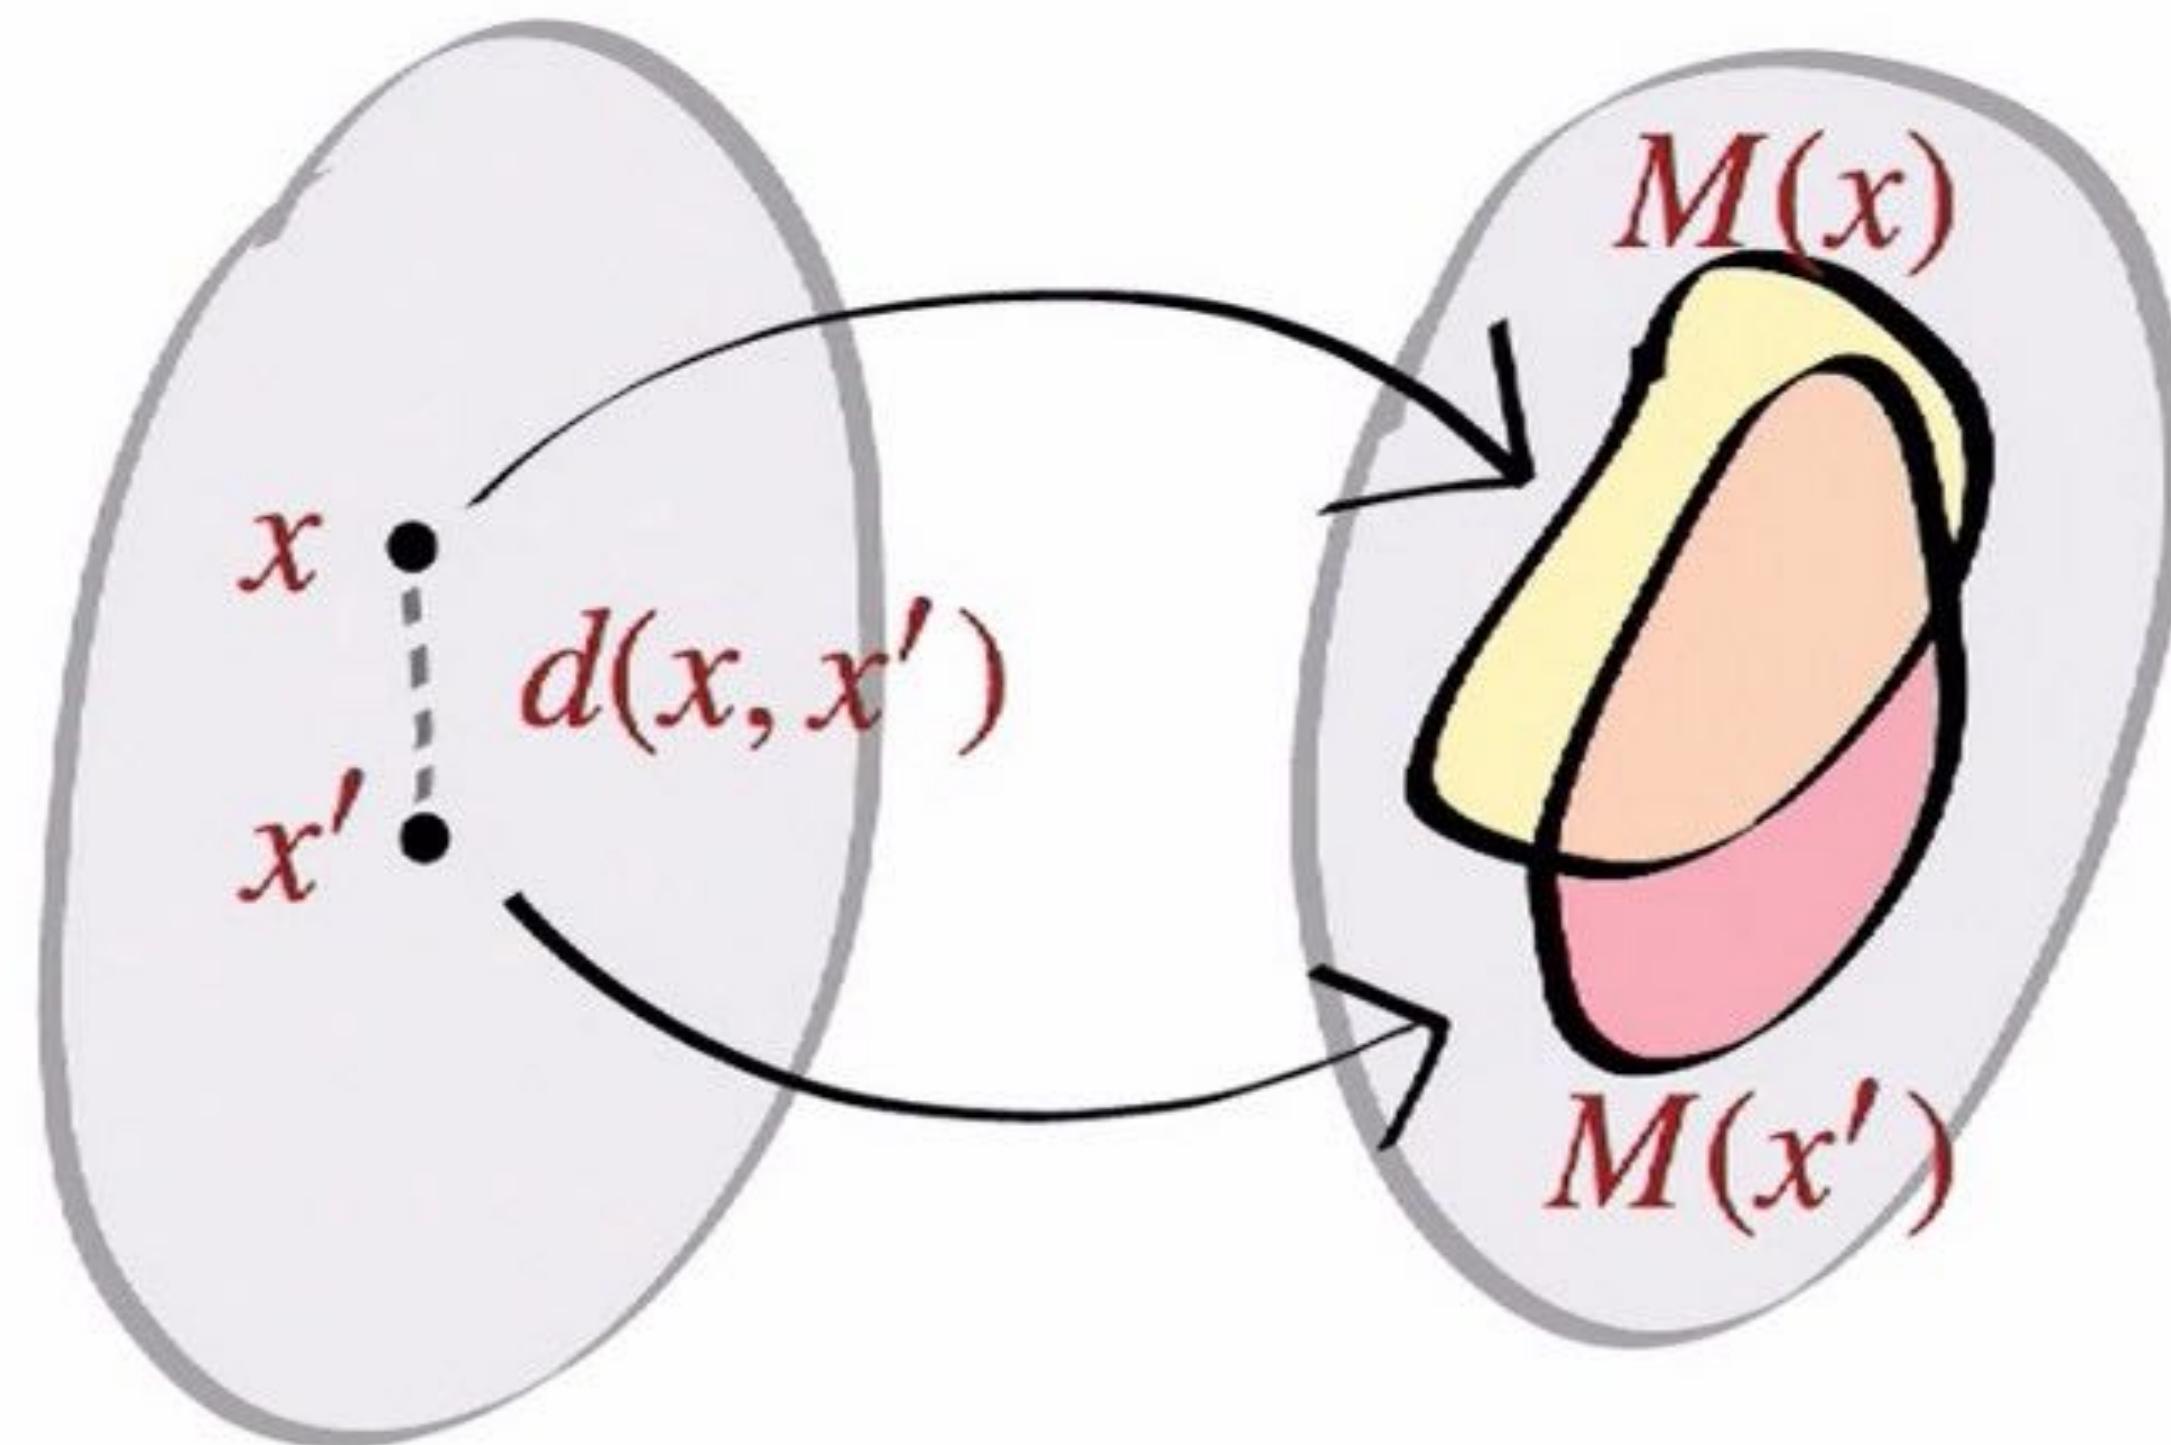


Similar distribution over outcomes

Formalize Individual Fairness

$M : x \rightarrow \Delta(O)$ Maps each individual example to a distribution of outcomes

$D(M(x), M(x')) \leq d(x, x')$ Where d and D are two distance functions





Fake Content

<https://www.youtube.com/watch?v=cQ54GDm1eL0>

Example 1: Fake Obama Video



can make it look like
anyone is saying anything

Example 2: Fake face Images by GAN

- Which are real/fake?



Example 2: Fake face Images by GAN

- Which are real/fake? <https://thispersondoesnotexist.com/>



Example 3: fiction Generated by GPT-3

- Completing a prompt from “Harry Potter and the Methods of Rationality”:

“... If there were any other monster that could defeat you as easily as that one, then you would have died of it long ago. That monster is stupidity. And that is why, my young apprentices, you must never, never, NEVER use the Killing Curse on anything without a brain!”

Professor Quirrell was now leaning on Harry’s desk.

Professor Quirrell stared straight into the eyes of every single student.

“The Killing Curse is too good for something without a brain. You will be fighting brains, or something near enough that makes no real difference. You will not be fighting trolls. You will not be fighting Dementors. The Killing Curse is no tool for anything less than the third most perfect killing machine in all Nature. If you are not prepared to use it against a mountain troll, then you are not prepared to use it at all.

Now. Pay attention to yourselves as I cast a simple spell. Listen to your own thoughts as I tell you how stupid you are.”

Professor Quirrell started pointing his wand at the ceiling.

...”

Detecting Fake Content

Fake photos/videos can have drawbacks.





Privacy

Example 1: Netflix Prize Competition

- Netflix Dataset: 480189 users x 17770 movies



	movie 1	movie 2	movie 3	movie 4	movie 5	movie 6
Tom	5	?	?	1	3	?
George	?	?	3	1	2	5
Susan	4	3	1	?	5	1
Beth	4	3	?	2	4	2

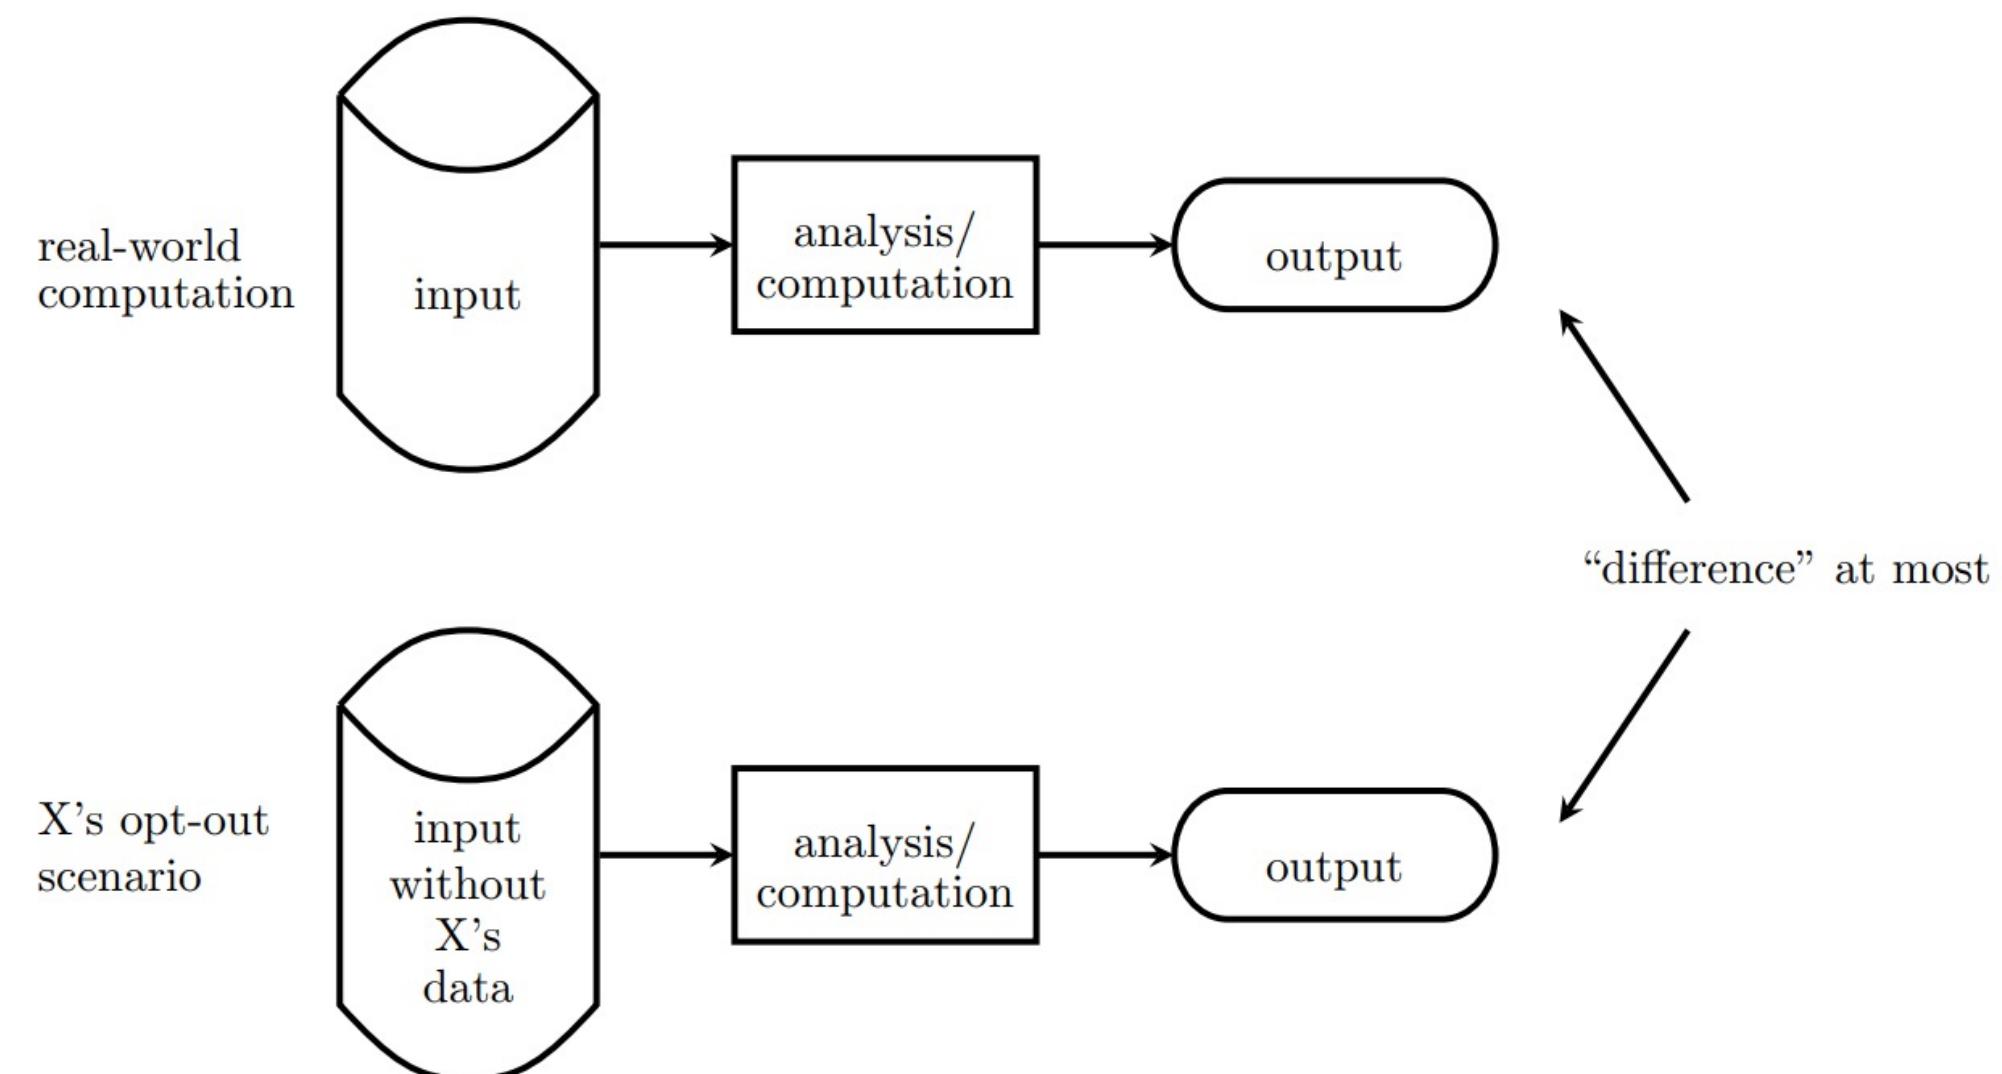
- The data was released by Netflix in 2006
 - replaced individual names with random numbers
 - moved around personal details, etc

Example 1: Netflix Prize Competition

- Arvind Narayanan and Vitaly Shmatikov compared the data with the non-anonymous IMDb users' movie ratings
- Very little information from the database was needed to identify the subscriber
 - simply knowing data about only two movies a user has reviewed allows for 68% re-identification success

Popular framework: Differential Privacy

- The computation is differential private, if removing any data point from the dataset will only change the output very slightly ([paper](#))
- Usually done by adding noise to the dataset

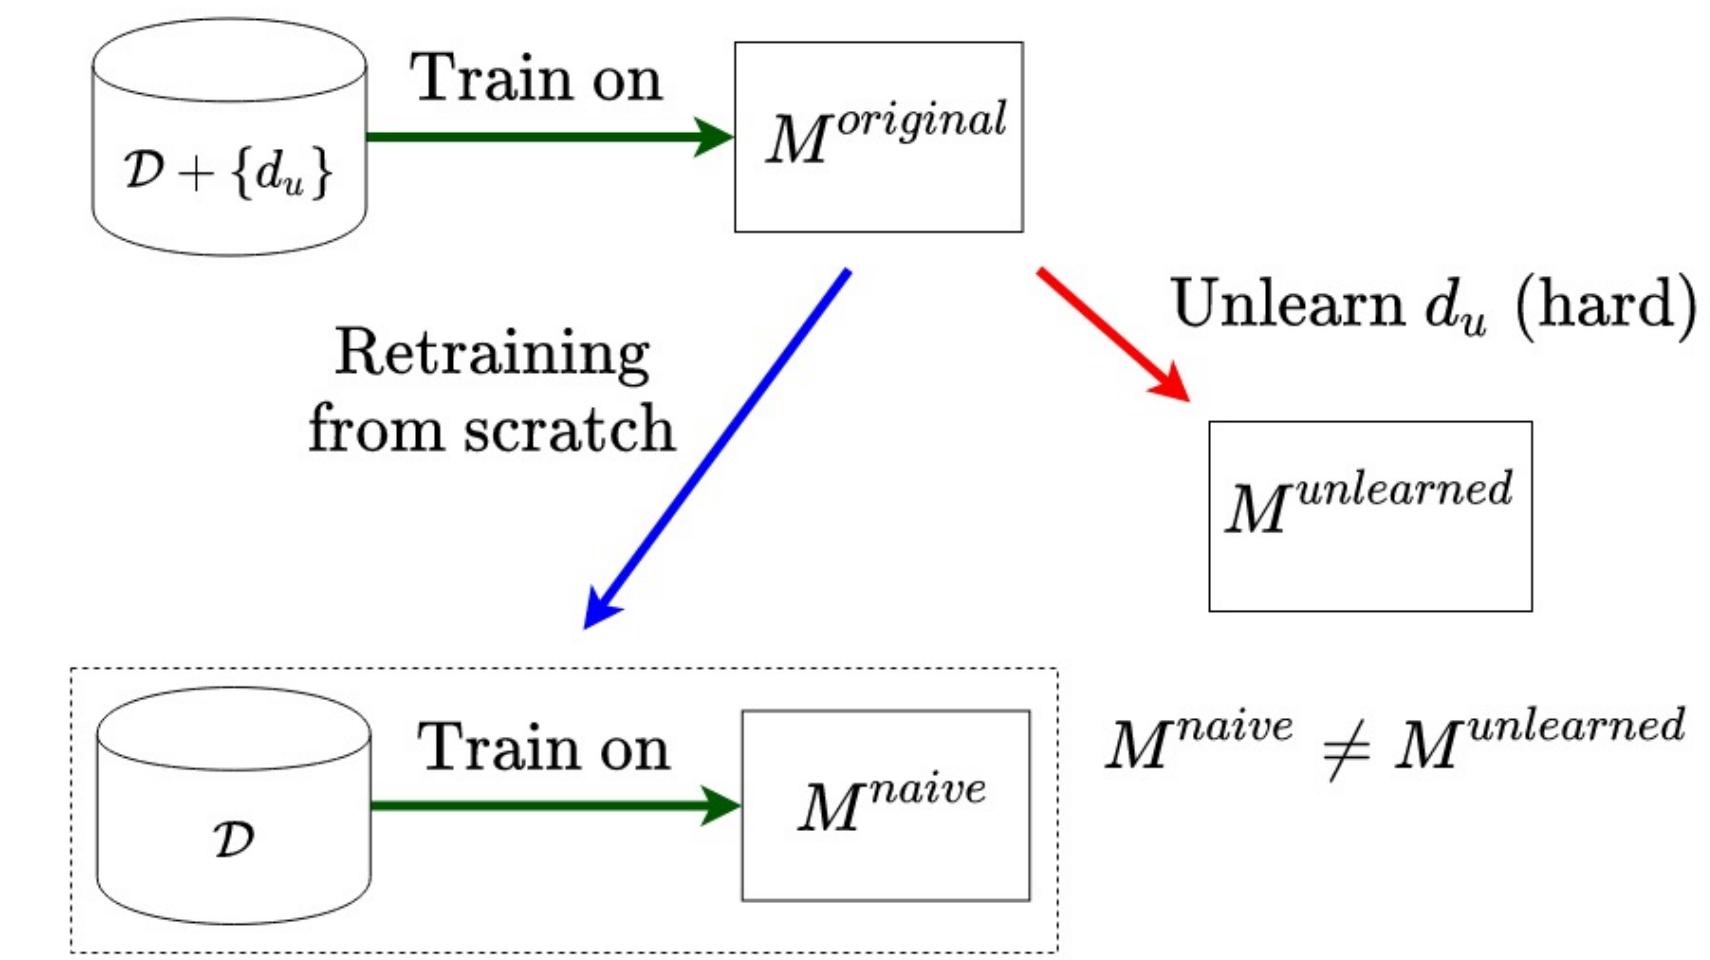


Right to be Forgotten

- The right to request that personally identifiable data be deleted
- E.g., an individual who did something foolish as a teenager doesn't want it to appear in web searches for the name for the rest of the life

Right to be Forgotten

- What if the data has been used in training a deep network?
 - Need to **unlearn**
- Other issues
 - Multiple copies of the data
 - Data already shared with others



From [Link](#)

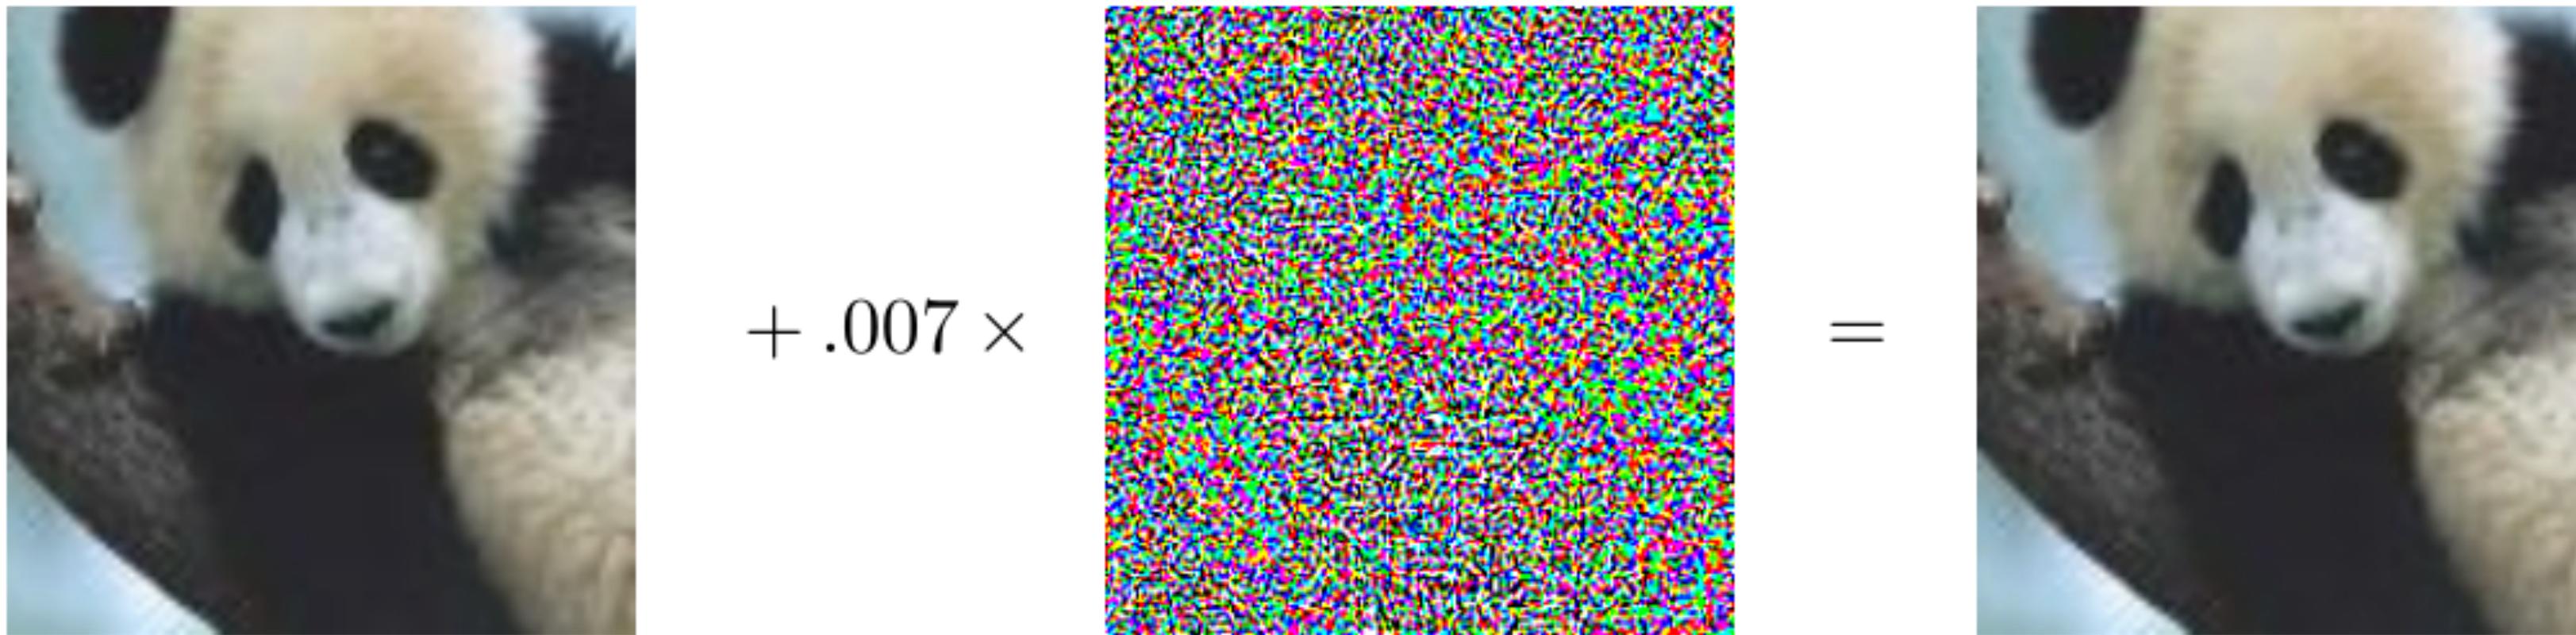


Adversarial Robustness

Adversarial Examples

“Inputs to ML models that an attacker has **intentionally** designed to cause the model to make a mistake”

Adversarial Examples



“Adversarial Classification” Dalvi et al 2004: fool spam filter

“Evasion Attacks Against Machine Learning at Test Time” Biggio 2013: fool neural nets

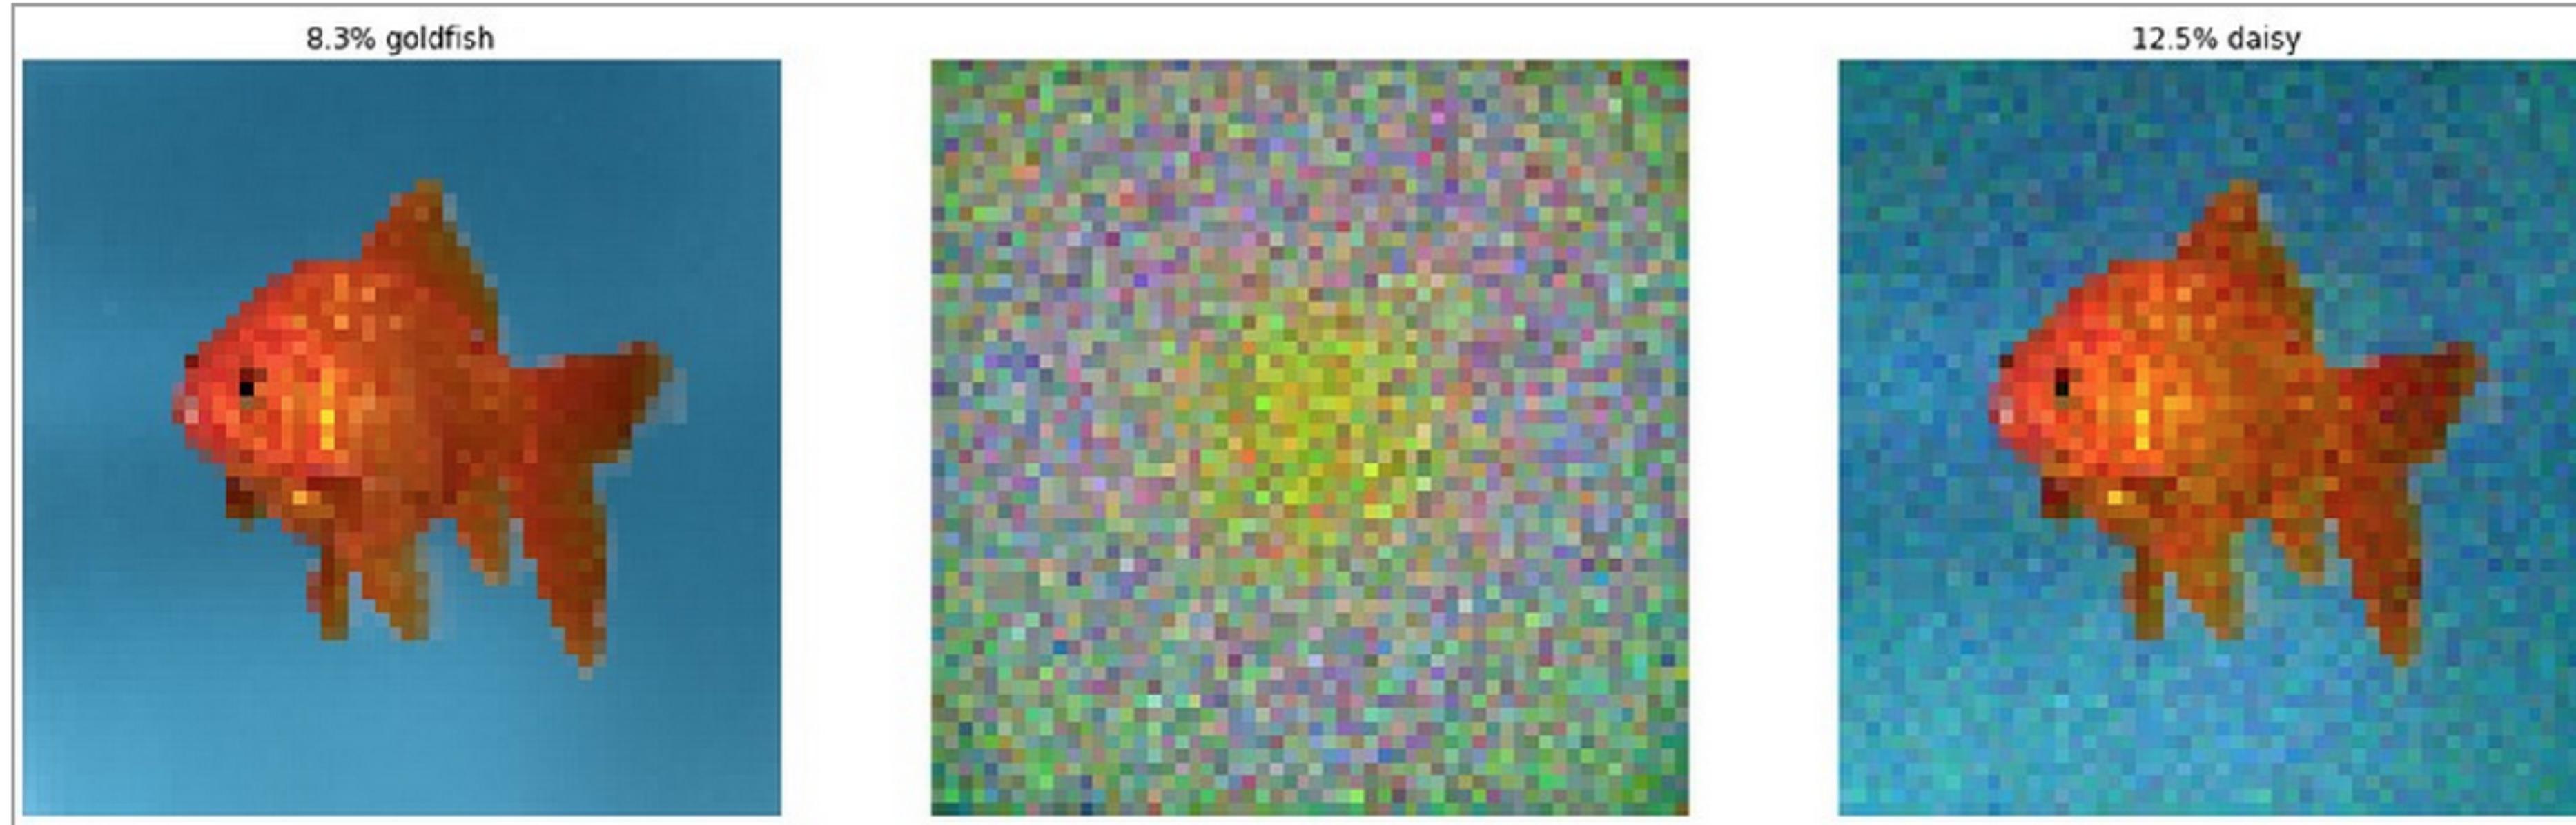
Szegedy et al 2013: fool ImageNet classifiers imperceptibly

Goodfellow et al 2014: cheap, closed form attack

Not just for neural nets

- Linear models
 - Logistic loss
 - Softmax loss
- Decision trees
- Nearest neighbors

Adversarial Examples Linear Models of ImageNet



(Andrej Karpathy, “Breaking Linear Classifiers on ImageNet”)

Adversarial Examples in NLP

Article: Super Bowl 50

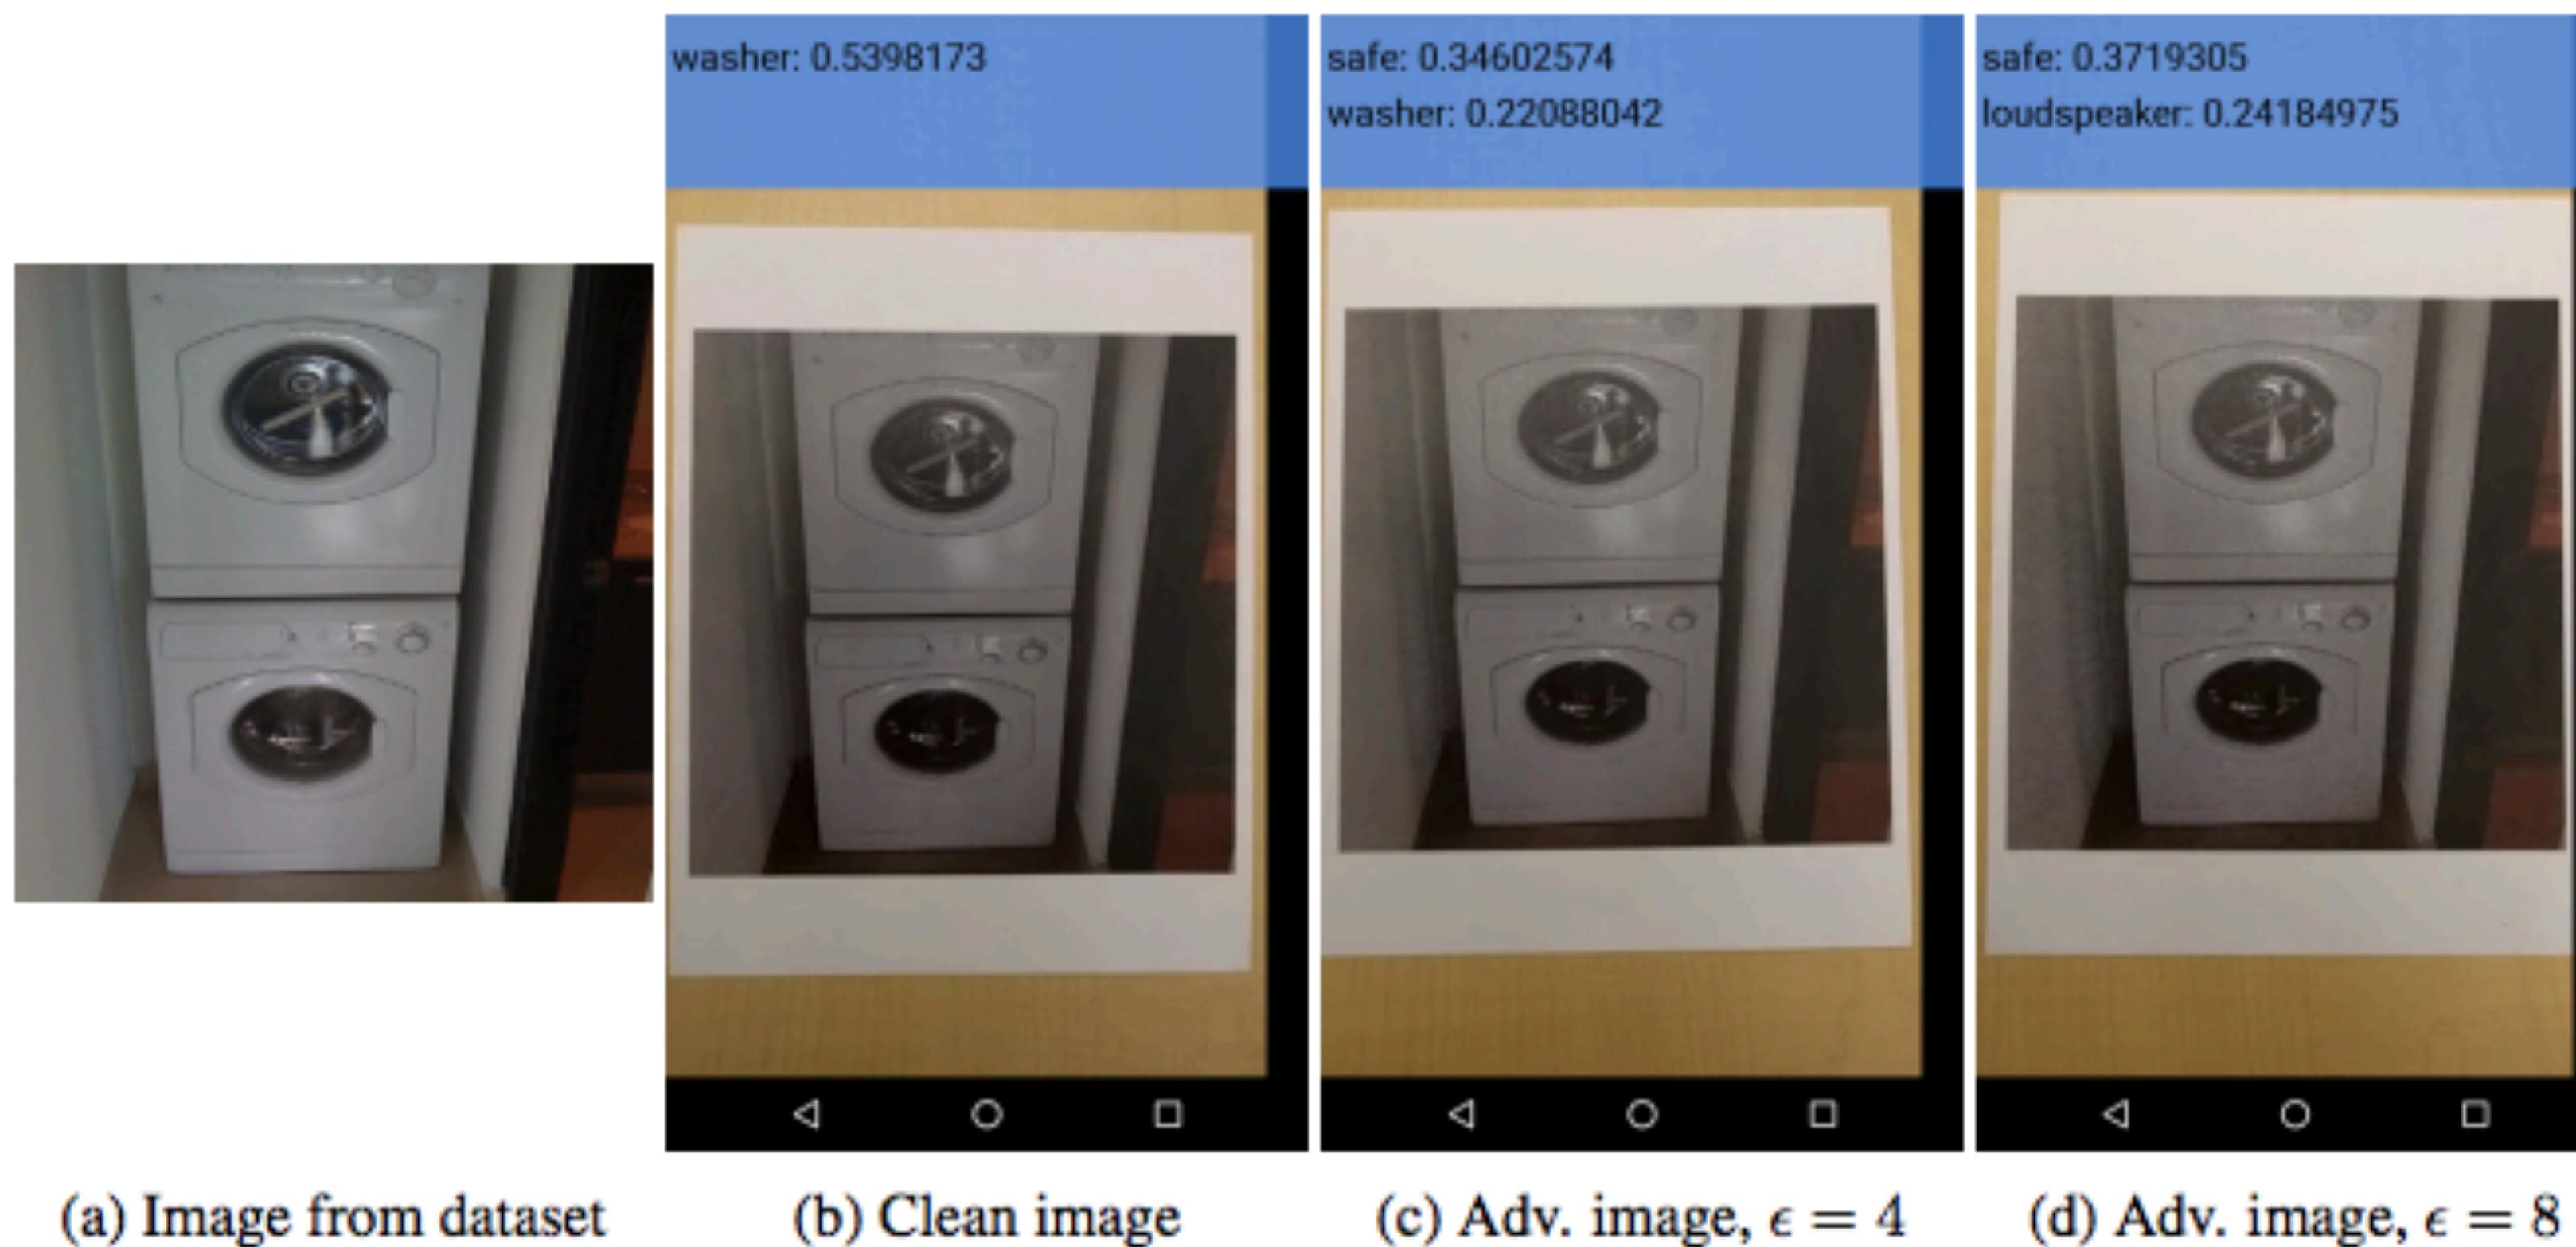
Paragraph: “*Peyton Manning became the first quarterback ever to lead two different teams to multiple Super Bowls. He is also the oldest quarterback ever to play in a Super Bowl at age 39. The past record was held by John Elway, who led the Broncos to victory in Super Bowl XXXIII at age 38 and is currently Denver’s Executive Vice President of Football Operations and General Manager. Quarterback Jeff Dean had jersey number 37 in Champ Bowl XXXIV.*

Question: “*What is the name of the quarterback who was 38 in Super Bowl XXXIII?*”

Original Prediction: John Elway

Prediction under adversary: Jeff Dean

Adversarial Examples in the Physical World



(Kurakin et al, 2016)

Adversarial Examples in the Physical World

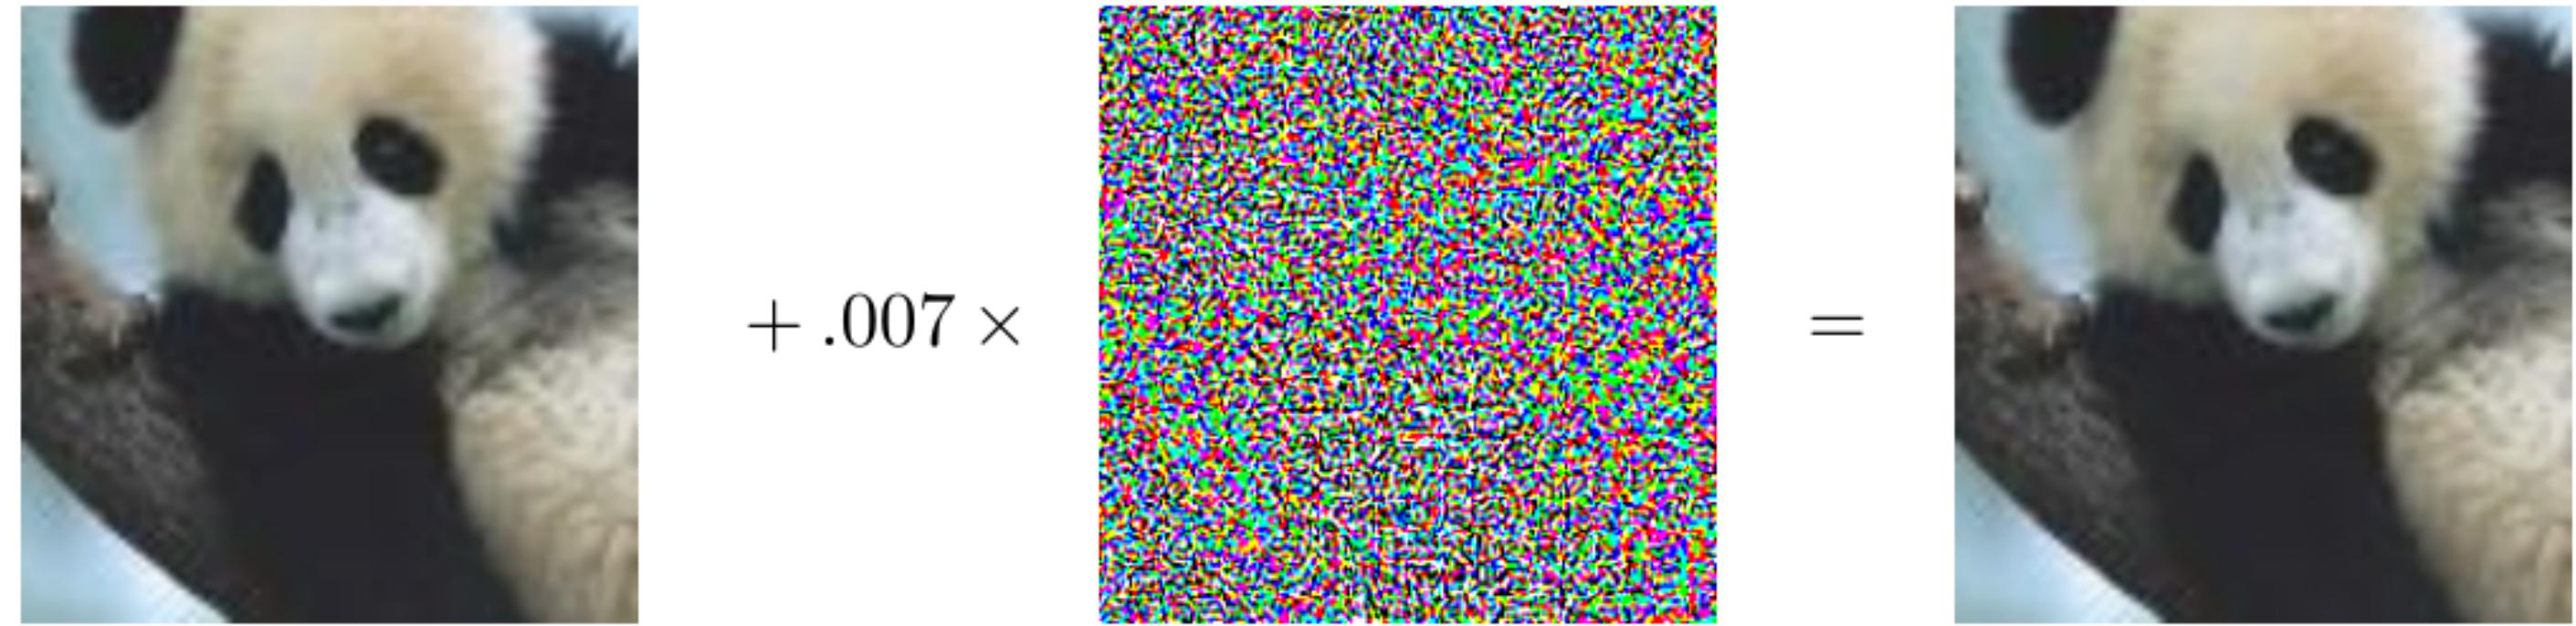


Fooling Image Recognition with Adversarial Examples

https://www.youtube.com/watch?v=piYnd_wY1T8

Generating Adversarial Examples

Simple approach: Fast Gradient Sign Method (FGSM) [Goodfellow et. al 2014]



\mathbf{x}
“panda”
57.7% confidence

$\text{sign}(\nabla_{\mathbf{x}} J(\boldsymbol{\theta}, \mathbf{x}, y))$
“nematode”
8.2% confidence

$\mathbf{x} + \epsilon \text{sign}(\nabla_{\mathbf{x}} J(\boldsymbol{\theta}, \mathbf{x}, y))$
“gibbon”
99.3 % confidence

$$\|\tilde{\mathbf{x}} - \mathbf{x}\|_\infty \leq \epsilon$$

$$\Rightarrow \tilde{\mathbf{x}} = \mathbf{x} + \epsilon \text{sign} (\nabla_{\mathbf{x}} J(\mathbf{x})) .$$

Defense: Adversarial Training

Labeled as bird



Still has same label (bird)

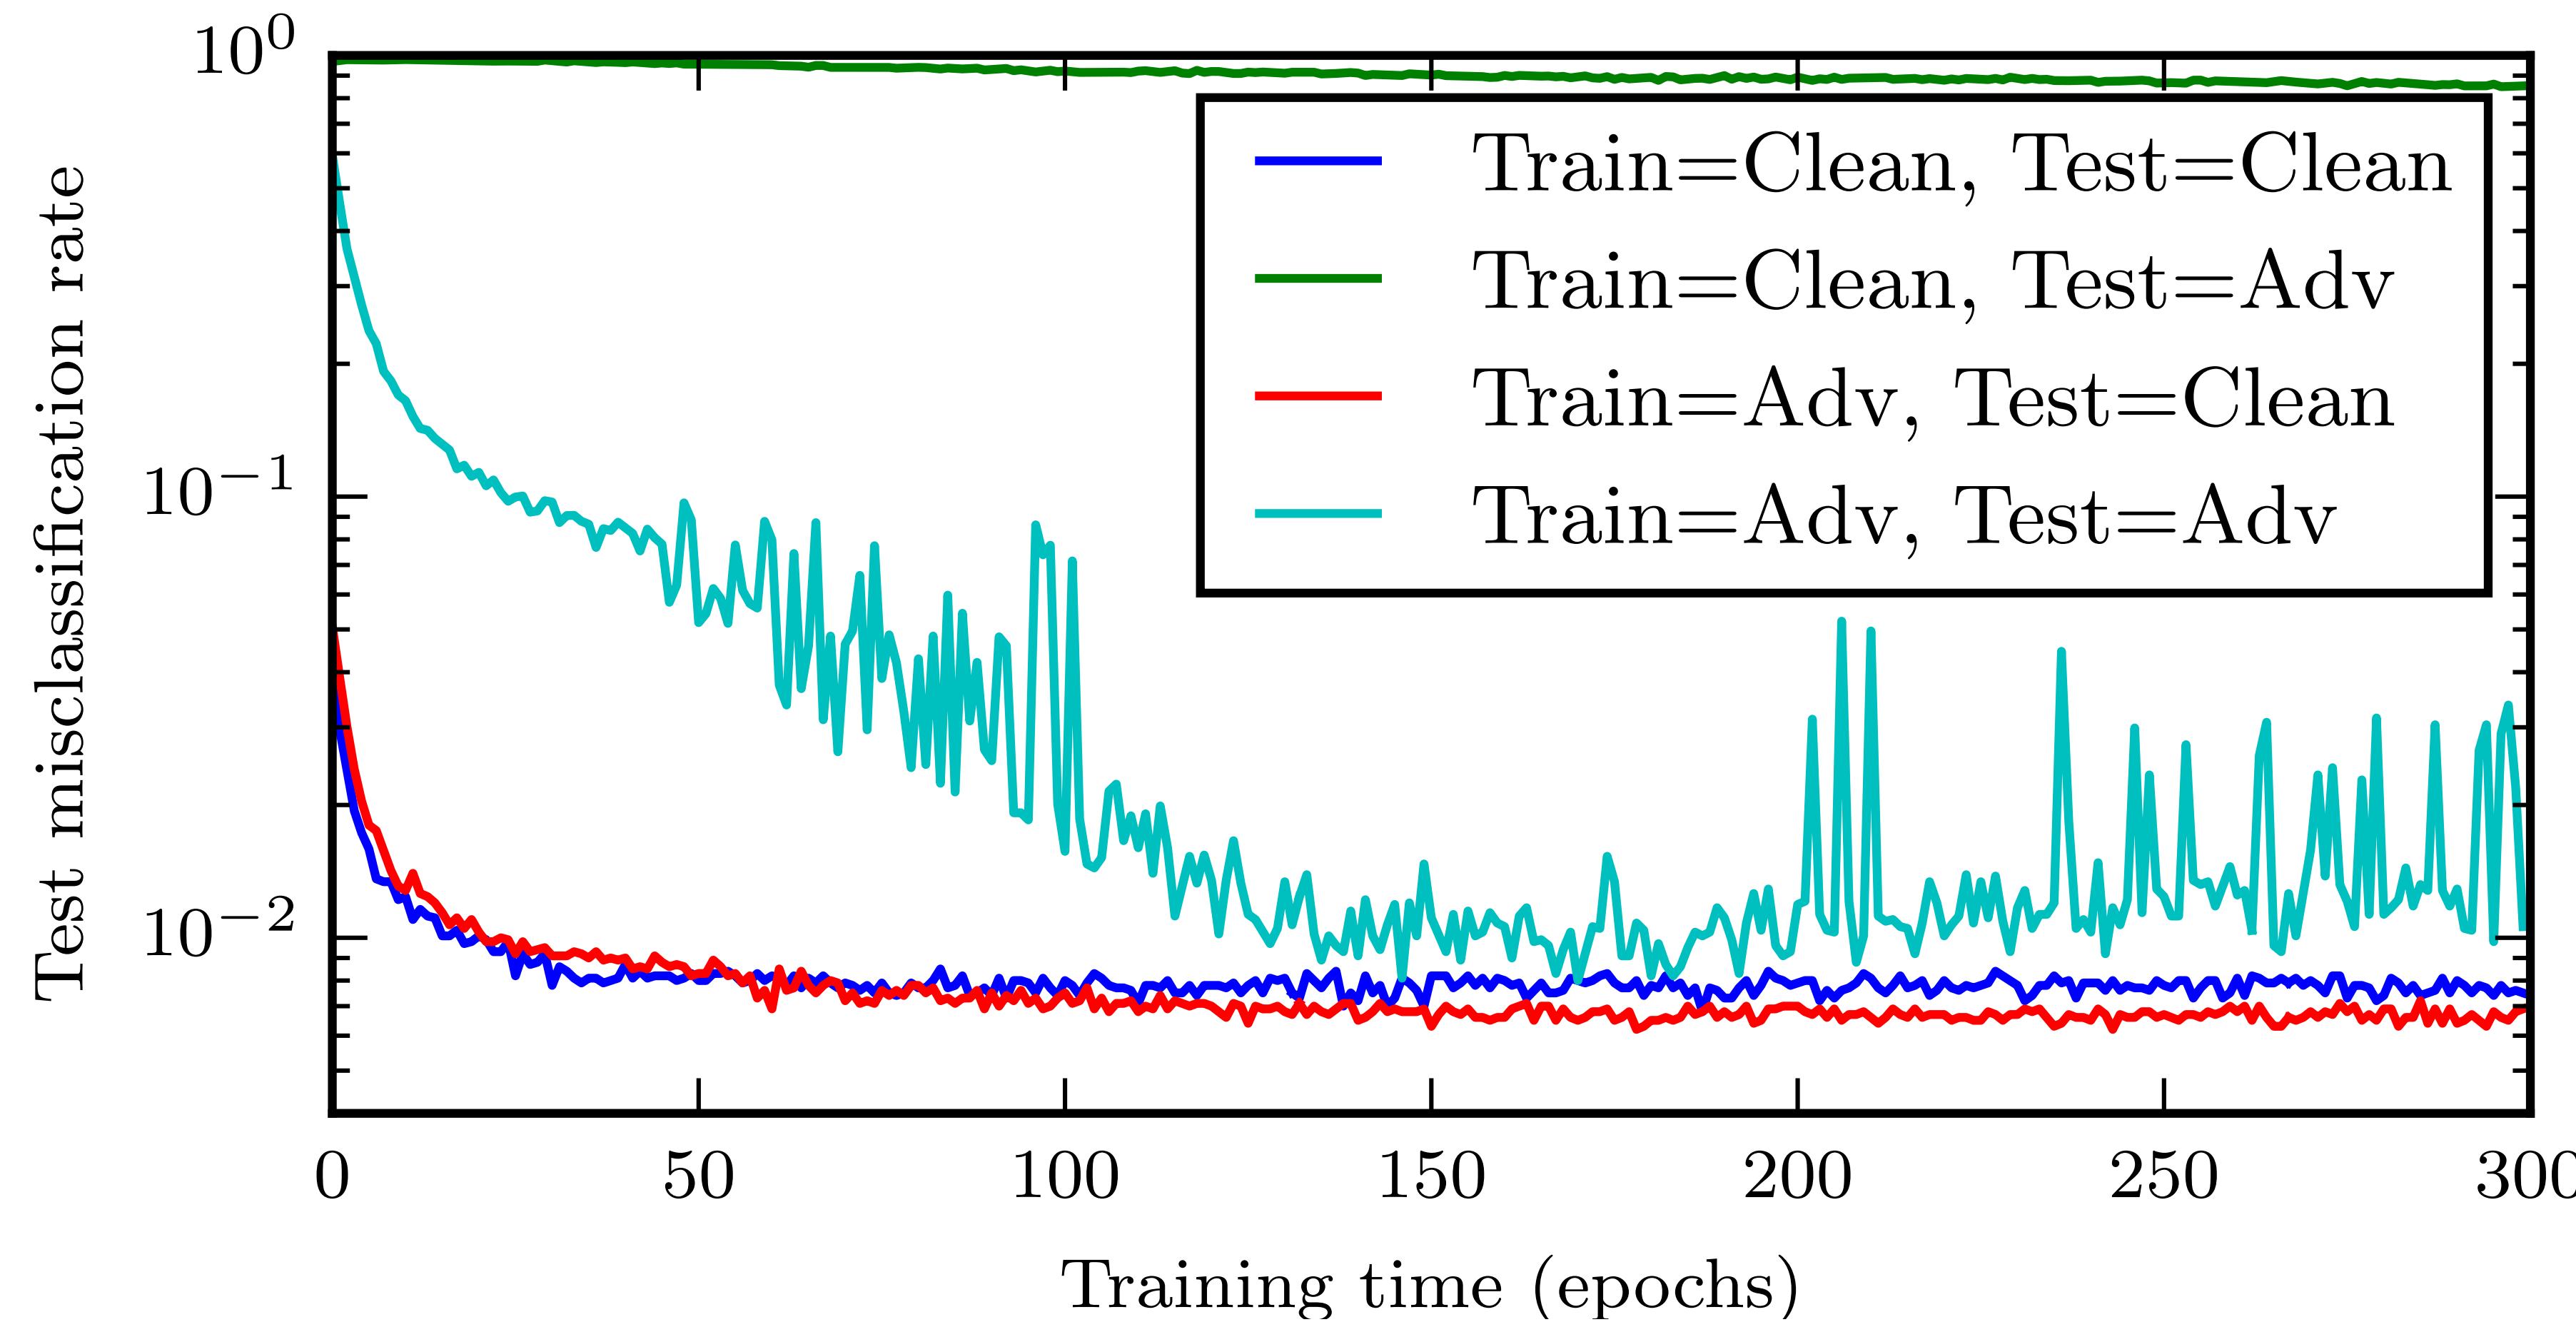


Decrease
probability
of bird class

Defense: Adversarial Training

Adversarial training can be viewed as **augmenting** the training data with adversarial examples.

Training on Adversarial Examples



Summary of Topics in Trustworthy ML

- Bias and Fairness
- Fake Content
- Adversarial robustness
- Privacy



Acknowledgement:

Some of the slides in these lectures have been adapted/borrowed from materials developed by Anthony Glitter, Yingu Liang, Hanxiao Liu: <http://www.cs.cmu.edu/~hanxiaol/slides/adversarial.pdf>, Ian Goodfellow